



Let's do it with complete mathematical precision. No shortcuts. We will use proper scaling ($\sqrt{d_k}$), true exponential softmax (e^x), and a standard ReLU activation ($\max(0, x)$).

To ensure the numbers don't explode into tedious fractions for manual calculation, I have carefully calibrated the input and weight matrices so that the exponents work out to clean integers (e^1, e^2 , etc.).

The Challenge Sequence

- **Prompt:** "Go home" (2 tokens)
- **Vocabulary:** ["now", "yesterday"] ($V = 2$)
- **Hidden Dimension:** $d = 2$

Your Setup Data

1. Input Matrix (X)

- $x_1 = \text{"Go"} = \begin{bmatrix} 1 & 1 \end{bmatrix}$
- $x_2 = \text{"home"} = \begin{bmatrix} 2 & 1 \end{bmatrix}$

$$X = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix}$$

2. Self-Attention Weights

$$W_Q = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad W_K = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad W_V = \begin{bmatrix} 1 & 2 \\ 2 & 0 \end{bmatrix}$$

- *Note:* Since $d_k = 2$, your scaling factor is exactly $\sqrt{d_k} = \sqrt{2} \approx 1.414$.

3. Feed-Forward Network Weights

The FFN will project up to dimension 3, apply ReLU, and project back to dimension 2.

$$W_1 = \begin{bmatrix} 1 & 0 & -2 \\ 0 & 1 & 1 \end{bmatrix}, \quad b_1 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

$$W_2 = \begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 0 & 2 \end{bmatrix}$$

4. LM Prediction Head (W_{vocab})

$$W_{\text{vocab}} = \begin{bmatrix} 2 & 1 \\ 1 & -2 \end{bmatrix} \begin{array}{l} \leftarrow \text{weights for "now"} \\ \leftarrow \text{weights for "yesterday"} \end{array}$$

Your Mathematical Roadmap

1. **Compute Q, K, V :** $Q = XW_Q, K = XW_K, V = XW_V$.
2. **Attention Scores:** Calculate QK^T .
3. **Scale:** Divide every value in the score matrix by $\sqrt{2}$ (use 1.414).
4. **True Softmax:** For each row, compute $e^{\text{scaled score}}$ and divide by the row sum to get the Attention Matrix A . (*Tip:* $e^1 \approx 2.72, e^2 \approx 7.39, e^3 \approx 20.09, e^4 \approx 54.60$).
5. **Context Vector (Z):** Compute $Z = AV$.
6. **FFN Block:** For the last token ("home"):
 - Calculate $\mathbf{h}_1 = (\mathbf{z}_{\text{home}} \cdot W_1) + b_1$
 - Apply $\mathbf{h}_{\text{act}} = \max(0, \mathbf{h}_1)$
 - Calculate $\mathbf{h}_{\text{final}} = \mathbf{h}_{\text{act}} \cdot W_2$
7. **The LM Head Head-to-Head:** Multiply $\mathbf{h}_{\text{final}}$ by W_{vocab}^T to get the final logits. Apply a final softmax to see which word wins!

Grab a calculator, map out your matrices step-by-step, and drop your calculations below when you are ready to see if your model predicts "now" or "yesterday".

show the claculations



Here is the complete step-by-step mathematical execution of the network. Grab your calculator to follow along!

Step 1: Compute Q, K , and V

Since W_Q and W_K are identity matrices, Q and K remain exactly equal to the input matrix X :

$$Q = XW_Q = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix} \begin{array}{l} \leftarrow Q_{\text{Go}} \\ \leftarrow Q_{\text{home}} \end{array}$$

$$K = XW_K = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix} \begin{array}{l} \leftarrow K_{\text{Go}} \\ \leftarrow K_{\text{home}} \end{array}$$

For V :

$$V = XW_V = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 2 & 0 \end{bmatrix}$$

- **Row 1 ("Go"):**
 - $1(1) + 1(2) = 3$
 - $1(2) + 1(0) = 2 \rightarrow \mathbf{v}_1 = \begin{bmatrix} 3 & 2 \end{bmatrix}$
- **Row 2 ("home"):**
 - $2(1) + 1(2) = 4$
 - $2(2) + 1(0) = 4 \rightarrow \mathbf{v}_2 = \begin{bmatrix} 4 & 4 \end{bmatrix}$

$$V = \begin{bmatrix} 3 & 2 \\ 4 & 4 \end{bmatrix}$$

Step 2: Calculate Raw Attention Scores (QK^T)

We multiply our Query matrix by the transpose of the Key matrix ($K^T = \begin{bmatrix} 1 & 2 \\ 1 & 1 \end{bmatrix}$):

$$\text{Raw Scores} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} (1 \cdot 1 + 1 \cdot 1) & (1 \cdot 2 + 1 \cdot 1) \\ (2 \cdot 1 + 1 \cdot 1) & (2 \cdot 2 + 1 \cdot 1) \end{bmatrix} = \begin{bmatrix} 2 & 3 \\ 3 & 5 \end{bmatrix}$$

Step 3: Scale by $\sqrt{d_k}$

Our hidden key dimension is $d_k = 2$, so we scale by dividing every element by $\sqrt{2} \approx 1.414$:

$$\text{Scaled Scores} = \begin{bmatrix} 2/1.414 & 3/1.414 \\ 3/1.414 & 5/1.414 \end{bmatrix} \approx \begin{bmatrix} 1.414 & 2.122 \\ 2.122 & 3.536 \end{bmatrix}$$

Step 4: True Exponential Softmax

We apply the softmax function $\frac{e^{s_i}}{\sum e^j}$ to each row to get the attention weights.

Row 1 ("Go"):

- $e^{1.414} \approx 4.112$
- $e^{2.122} \approx 8.348$
- Sum = $4.112 + 8.348 = 12.460$
- Weights = $\begin{bmatrix} \frac{4.112}{12.460} & \frac{8.348}{12.460} \end{bmatrix} = \begin{bmatrix} 0.33 & 0.67 \end{bmatrix}$

Row 2 ("home"):

- $e^{2.122} \approx 8.348$
- $e^{3.536} \approx 34.329$
- Sum = $8.348 + 34.329 = 42.677$
- Weights = $\begin{bmatrix} \frac{8.348}{42.677} & \frac{34.329}{42.677} \end{bmatrix} = \begin{bmatrix} 0.20 & 0.80 \end{bmatrix}$

$$A = \begin{bmatrix} 0.33 & 0.67 \\ 0.20 & 0.80 \end{bmatrix}$$

Step 5: Compute Context Vector Matrix ($Z = AV$)

Now blend the normalized weights with our Value matrix V :

$$Z = \begin{bmatrix} 0.33 & 0.67 \\ 0.20 & 0.80 \end{bmatrix} \begin{bmatrix} 3 & 2 \\ 4 & 4 \end{bmatrix}$$

- **Row 1 (Z_{Go}):**
 - $0.33(3) + 0.67(4) = 0.99 + 2.68 = 3.67$
 - $0.33(2) + 0.67(4) = 0.66 + 2.68 = 3.34$
- **Row 2 (Z_{home}):**
 - $0.20(3) + 0.80(4) = 0.60 + 3.20 = 3.80$
 - $0.20(2) + 0.80(4) = 0.40 + 3.20 = 3.60$

$$Z = \begin{bmatrix} 3.67 & 3.34 \\ 3.80 & 3.60 \end{bmatrix}$$

Step 6: The Feed-Forward Network (FFN)

Because this is an autoregressive task predicting the next token, we only isolate and process the context vector of the **last token ("home")**:

$$\mathbf{z}_{\text{home}} = \begin{bmatrix} 3.80 & 3.60 \end{bmatrix}$$

1. First Projection Layer (W_1):

$$\mathbf{h}_1 = \mathbf{z}_{\text{home}} \cdot W_1 = \begin{bmatrix} 3.80 & 3.60 \end{bmatrix} \begin{bmatrix} 1 & 0 & -2 \\ 0 & 1 & 1 \end{bmatrix}$$

- Dim 1 = $3.80(1) + 3.60(0) = 3.80$
- Dim 2 = $3.80(0) + 3.60(1) = 3.60$
- Dim 3 = $3.80(-2) + 3.60(1) = -7.60 + 3.60 = -4.00$

$$\mathbf{h}_1 = \begin{bmatrix} 3.80 & 3.60 & -4.00 \end{bmatrix}$$

2. Activation Function (ReLU):

$$\mathbf{h}_{\text{act}} = \max(0, \mathbf{h}_1) = \begin{bmatrix} 3.80 & 3.60 & 0.00 \end{bmatrix}$$

(Notice how the third negative feature dimension gets completely zeroed out).

3. Second Projection Layer (W_2):

$$\mathbf{h}_{\text{final}} = \mathbf{h}_{\text{act}} \cdot W_2 = \begin{bmatrix} 3.80 & 3.60 & 0.00 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 0 & 2 \end{bmatrix}$$

- Dim 1 = $3.80(1) + 3.60(1) + 0.00(0) = 7.40$
- Dim 2 = $3.80(0) + 3.60(1) + 0.00(2) = 3.60$

$$\mathbf{h}_{\text{final}} = \begin{bmatrix} 7.40 & 3.60 \end{bmatrix}$$

Step 7: The LM Head Prediction

We pass this fully processed hidden representation through our vocabulary map (W_{vocab}) to find the final text predictions:

$$\text{Logits} = \mathbf{h}_{\text{final}} \cdot W_{\text{vocab}}^T = \begin{bmatrix} 7.40 & 3.60 \end{bmatrix} \begin{bmatrix} 2 & 1 \\ 1 & -2 \end{bmatrix}$$

- **Logit("now"):** $7.40(2) + 3.60(1) = 14.80 + 3.60 = \mathbf{18.40}$
- **Logit("yesterday"):** $7.40(1) + 3.60(-2) = 7.40 - 7.20 = \mathbf{0.20}$

Final Softmax Validation

- $e^{18.40} \approx 97,949,311.5$
- $e^{0.20} \approx 1.22$

$$P(\text{"now"}) = \frac{97,949,311.5}{97,949,311.5 + 1.22} \approx \mathbf{99.9999\%}$$

$$P(\text{"yesterday"}) \approx \mathbf{0.0001\%}$$

The network confidently completes your sentence as: **"Go home now"**.