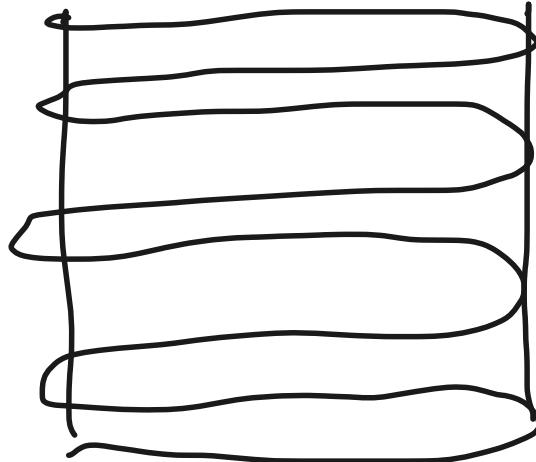


Simple Linear Rgression (SLR)
Multiple Linear Regression (MLR)
Polynomial Regression
Regression Outcomes & Interpretation

This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

Stats

ML



Levels of Measurement :

NOMINAL

- Red, Blue, Green
- Gender

height, weight

temperature

Gender

Grades

Salary

Price

Qty.

ORDINAL

- Grades

INTERVAL

- Temp.

Understanding

ML requires

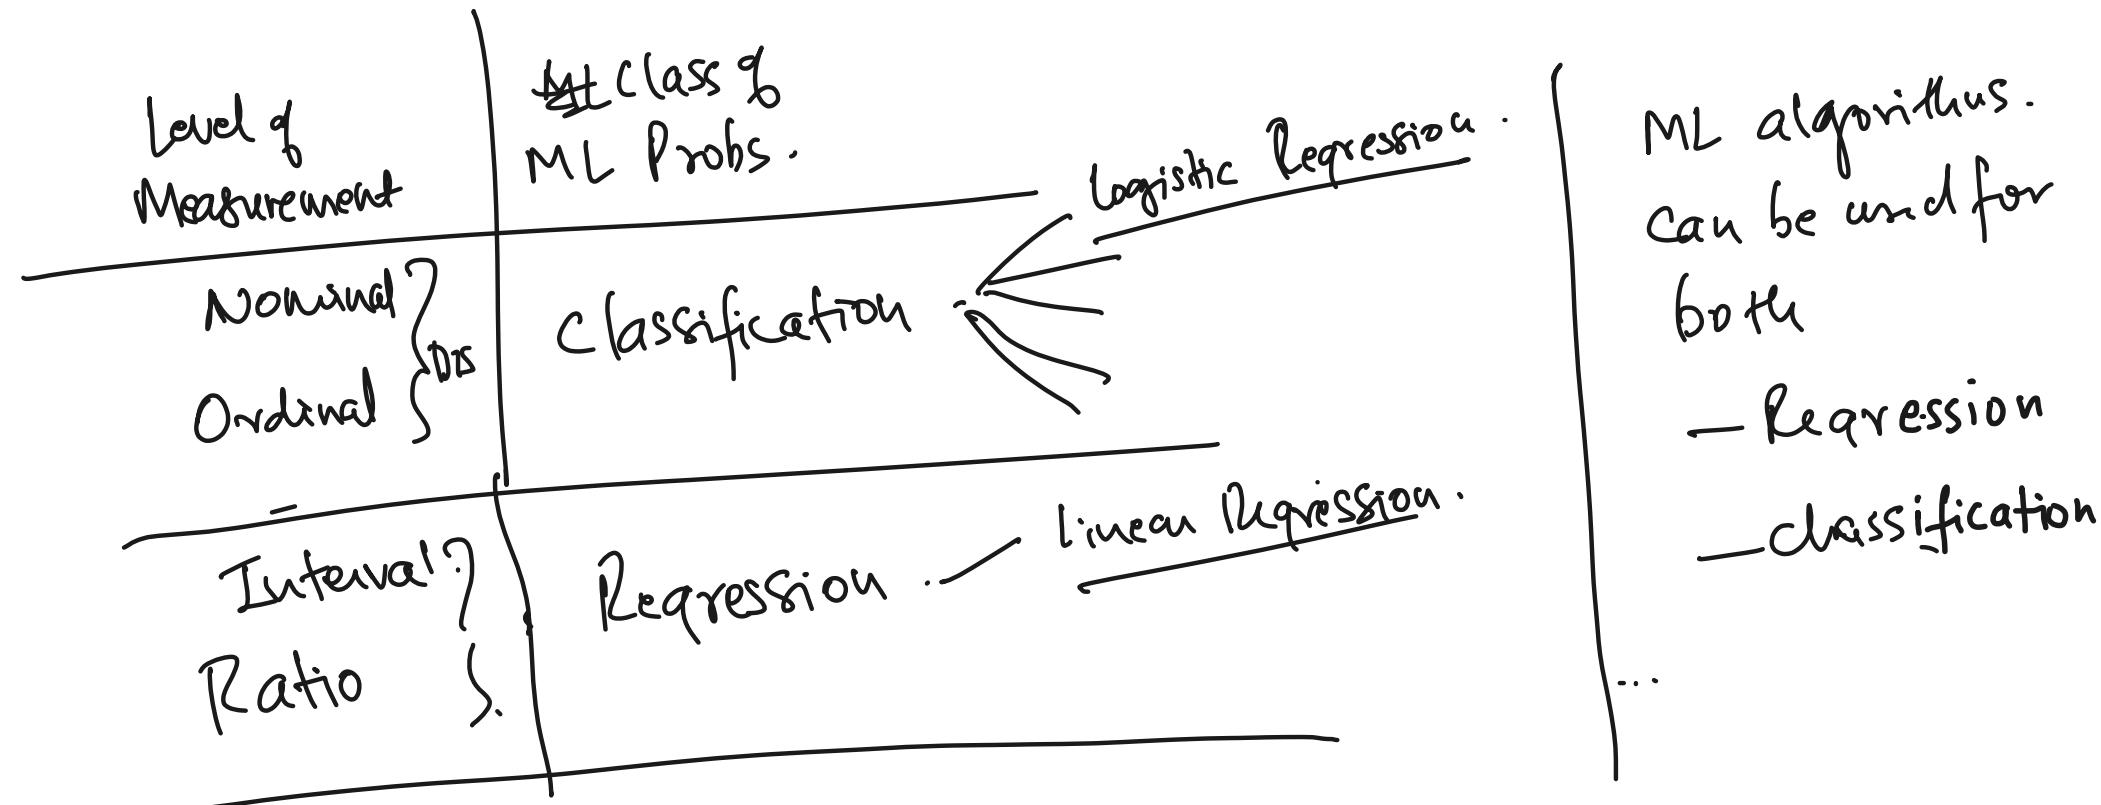
good understand.

- ding of Stats

RATIO - ht, wt

This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

Level of Measurement	Description	Typical Statistics/Parameters	Supported Operations
Nominal	Categorical data without any order	<ul style="list-style-type: none"> - Frequency counts - Mode 	None
Ordinal	Categorical data with a meaningful order	<ul style="list-style-type: none"> - Frequency counts - Mode - Median 	<ul style="list-style-type: none"> - Comparison of order (e.g., $>$, $<$, $=$)
Interval	Numeric data with equal intervals but no true zero	<ul style="list-style-type: none"> - Frequency counts - Mode - Median - Mean - Standard deviation - Correlation (e.g., Pearson's) 	<ul style="list-style-type: none"> - Addition - Subtraction
Ratio	Numeric data with equal intervals and a true zero	<ul style="list-style-type: none"> - Frequency counts - Mode - Median - Mean - Standard deviation - Range - Ratio comparisons - Geometric mean - Harmonic mean 	<ul style="list-style-type: none"> - Addition - Subtraction - Multiplication - Division



$$y = f(x)$$

If y is available \rightarrow Supervised ML.

If y is unavailable \rightarrow

Unsupervised ML.

clustering

Regression
classification

More
discovery
necessary.

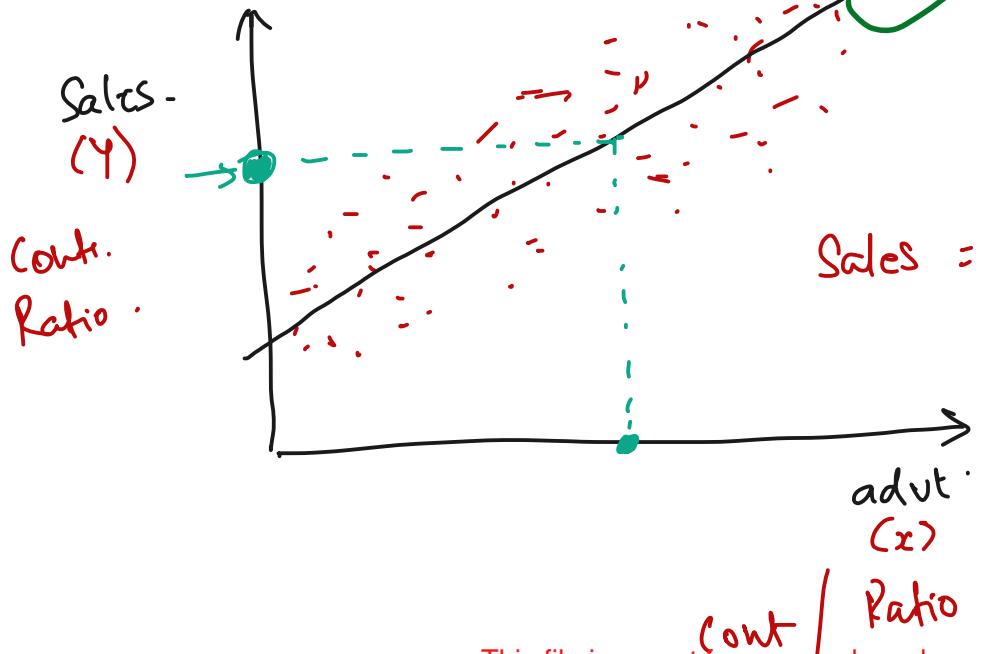
ML

Predicted | Response | output
Inputs | features | predictors

$$y = f(x)$$

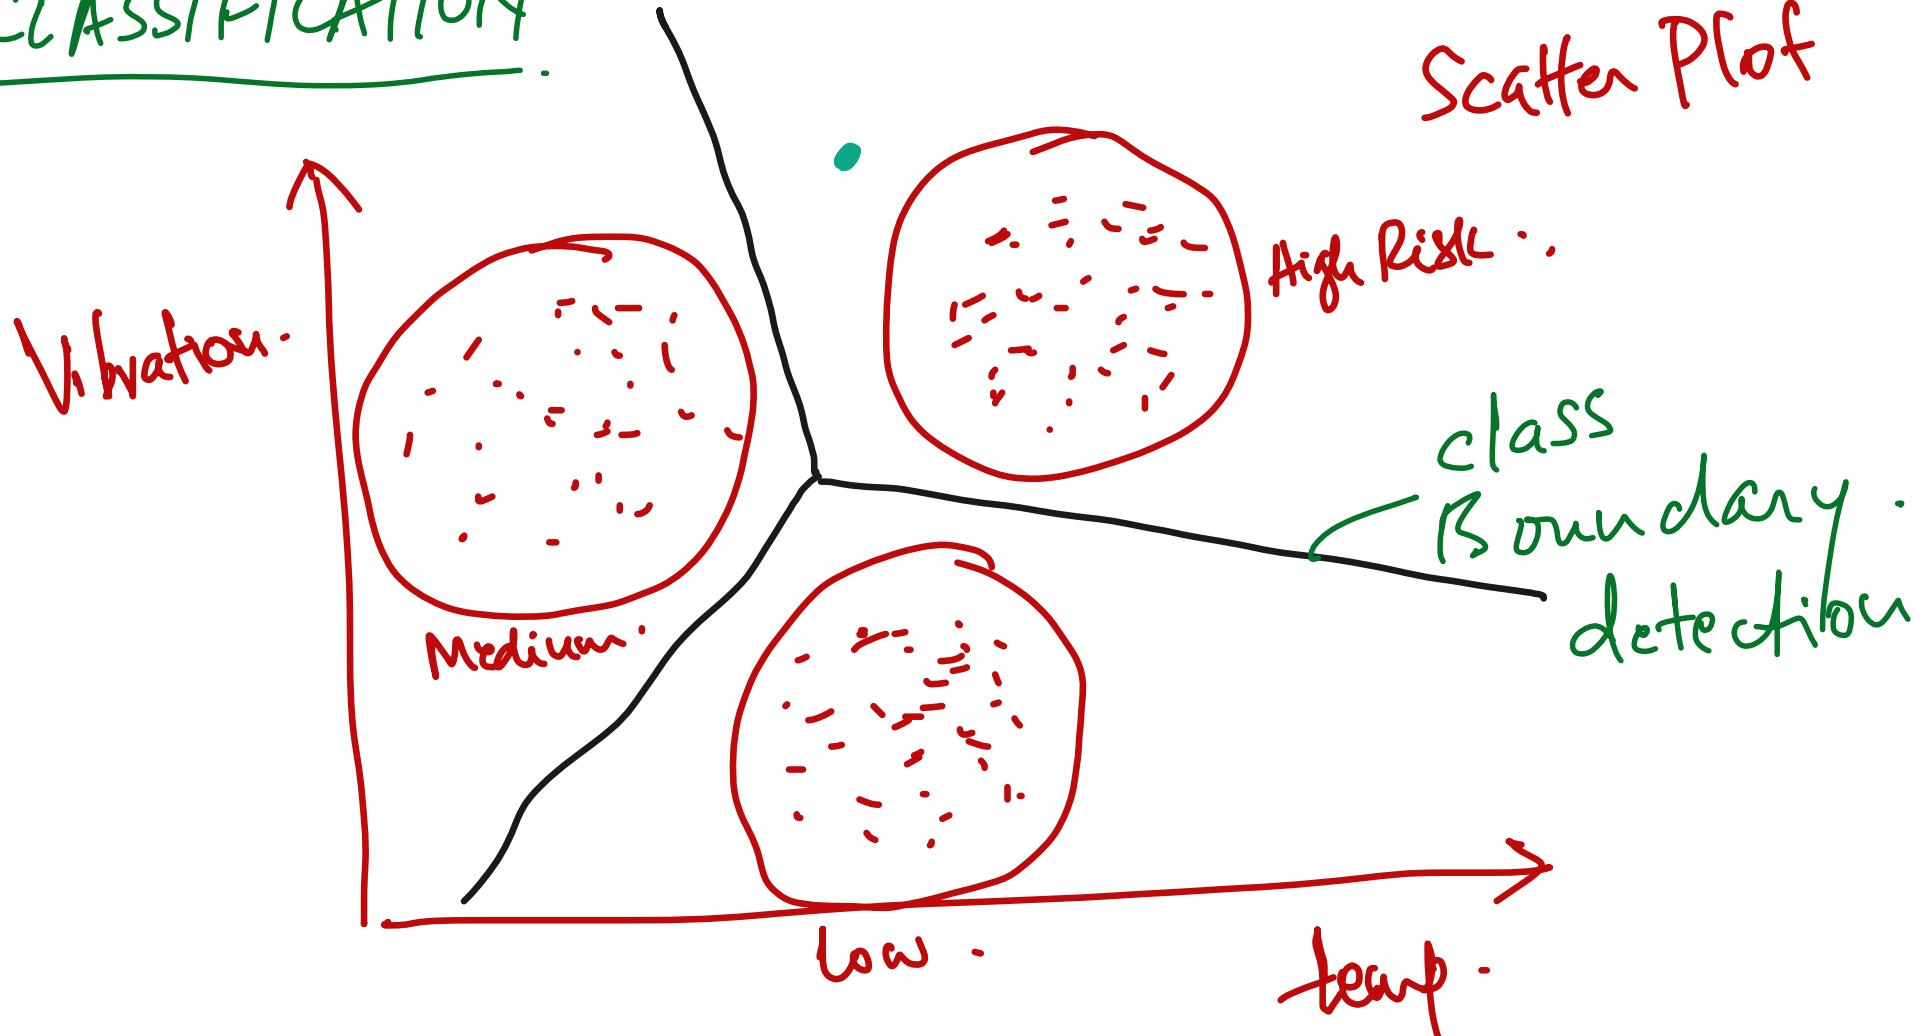
REGRESSION.

Prediction 'line' / curve



Sales = $f(advt, \text{tf}, \text{earnings}, \text{edu}, \text{weather}, \text{connectivity})$.
[$x_1, x_2, x_3, x_4, \dots$] $\rightarrow [+]$.

CLASSIFICATION



Exercise - 1

- In your spreadsheet explore how to enable the 'Data Analysis' toolpak, or equivalent, to perform statistical calculations!
- Use the uploaded data set data-set-for-SLR-2025.csv to perform Simple Linear Regression and generate the output (for this exercise, use the entire data for training the model)
- Analyze the output generated
- Now, using the **LinearRegression** (sklearn) function of Python create a regression model and calculate metrics like R2, MAE, RMSE and analyze the results
- Further, use the **OLS** function from 'statsmodels' package to perform regression, print the results and review them.

These steps need to be completed before proceeding!

	A	B	C	D	E	F	G	H	I	J	K	L
1	y	x				SUMMARY OUTPUT						
2	7.238462	0.025641				Regression Statistics						
3	6.310256	0.051282	OUTPUTS CREATED			Multiple R	0.906270151					
4	8.315385	0.076923	By REGRESSION			R Square (R^2)	0.821325586					
5	4.787179	0.102564	TOOLS / FUNCTIONS			Adjusted R Square	0.819483582					
6	5.592308	0.128205				Standard Error	1.882513522					
7	7.830769	0.153846				Observations	99					
8	9.902564	0.179487										
9	5.607692	0.205128										
10	5.146154	0.230769	EXCEL (or									
11	4.784615	0.25641										
12	7.05641	0.282051	any other									
13	9.394872	0.307692	spreadsheet)									
14	6.8	0.333333										
15	4.871795	0.358974										
16	5.376923	0.384615										
17	10.71538	0.410256	'b'			Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	
18	11.55385	0.435897	'a'			Intercept	5.922586409	0.381284372	15.53325	4.65E-28	5.165842474	6.679330343
19	9.258974	0.461538				x	5.452241187	0.258203842	21.11603	4.72E-38	4.939778036	5.964704333
20	8.097436	0.487179										
21	12.10256	0.512821										

What is a good model?

- One that explains most of the variations in the data.

$$\sum (y_i - \bar{y})^2 = SST \quad (\text{SST} = \text{measure of total variation in the given dataset})$$

$$\sum (y_i - \hat{y}_i + \hat{y}_i - \bar{y})^2$$

$$\sum \left[(y_i - \hat{y}_i)^2 + (\hat{y}_i - \bar{y})^2 + 2(y_i - \hat{y}_i) \cdot (\hat{y}_i - \bar{y}) \right] =$$
$$\sum (y_i - \hat{y}_i)^2 + \sum (\hat{y}_i - \bar{y})^2 + 2 \sum (y_i - \hat{y}_i) \cdot (\hat{y}_i - \bar{y})$$

$$SST = SSE + SSR + 2 \sum (y_i - \hat{y}_i) \cdot (\hat{y}_i - \bar{y})$$

SST => total variation explained by the regression model

SSE => variation NOT explained by the model, attributed to random errors

$$SST = SSR + SSE + \text{ZERO}$$

$$1 = \frac{SSR}{SST} + \frac{SSE}{SST}$$

$$1 = R^2 + \frac{SSE}{SST}$$

$$R^2 = 1 - \frac{SSE}{SST}$$

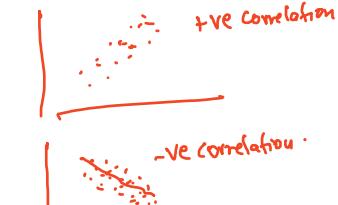
R^2 is the square of the correlation coefficient ' r '! } S.L.R

R^2 = COEFFICIENT OF DETERMINATION (C.O.D.)

= square of the correlation coefficient ' r ' between x & y .

SELF STUDY

- CORRELATION
- CORRELATION COEFF



We usually calculate SSE, MSE, RMSE, MAE, R2 as metrics reflecting the quality of Linear Regression. However, when we use built-in LR functionality, in tools like Excel, many more numbers are generated .. as shown below. What are they and how to interpret / use them?

ANOVA						
	df	SS	MS	F	Significance F	
Regression	1	1580.159507	1580.16	445.8869	4.72338E-38	
Residual	97	343.7541446	3.543857			
Total	98	1923.913652				

	Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%
Intercept (b)	5.922586409	0.381284372	15.53325	4.65E-28	5.165842474	6.679330343
x (a)	5.452241187	0.258203842	21.11603	4.72E-38	4.939778036	5.964704338

Regression Statistics	
Multiple R	0.906270151
R Square	0.821325586
Adjusted R Square	0.819483582
Standard Error	1.882513522
Observations	99

$$y = b + a x \quad \text{or}$$

$$y = \beta_0 + \beta_1 x$$

- To understand these numbers we have to go back to the basics of statistics.
- We need to start with the fact that the (y,x) data that we have is essentially a **sample** (in this case 1 sample of 99 observations)
- We have fitted an LR model using this sample. Therefore the calculated values of **a** and **b** are only an estimate of the population's **actual** a and b.
- Our aim, really, is to predict the value of **y** for an **x** that is not a part of the sample. That is, we need a model that is '**general**' and which reflects the reality of the population, and not limited to the sample that we have.
- So we really need to know **how good** an estimate these calculated values (a, b) are. Are they really usable? How much confidence should we have on our calculations?
- This is where we need to understand the concepts, from statistics, of **sampling distributions** and **confidence intervals**

We conduct some 'thought' experiments, related to estimating the population mean from the sample mean:

- Assume that from a population we can take multiple **good, representative** samples, let's say k samples, each of size n . Let's call each sample as s_i
- Using each s_i , we calculate its mean and call it m_i
- Since our samples are **good, representative** samples of the population, they will result in means m_i that are close to each other (why ? try to reason this out)
- If we collect all the m_i and create a frequency table and a histogram, its shape will be as shown below.



- We will observe that such a histogram indicates that the calculated means m_i tend to have Normal Distribution (as per the **Central Limit Theorem** - see next slide)
- This distribution is known as the **Sampling Distribution of the mean** or **Sampling Distribution of the sample mean** and it has the following properties:
 - The **Expected Value** (ie. mean) of such a distribution is very close to the population mean
 - The Standard Deviation of this distribution - known as the **Standard Error**, and denoted by $S_{\bar{x}}$ - is related to **sigma**, the population's standard deviation in the following way:

$$S_{\bar{x}} = \frac{\sigma}{\sqrt{n}} \quad \text{where } n = \text{Size of the Sample} .$$

- Implication of this formula: For a given population, with a given sigma, $S_{\bar{x}}$ reduces with increase in the sample size **n**. This, in turn, indicates less uncertainty in estimating the true value of the population mean.
- This appeals to our common sense that **as the sample sizes increase, our analysis becomes more accurate** or, conversely, **smaller sample sizes result in more uncertainty or inaccuracy in our predicted results**

So - given 100 observations, ~~does it make sense to say it is as 1 sample of size 100, or 10 samples of size 10?~~
 Sharing or publishing the contents in part or full is liable for legal action.

The Central Limit Theorem

The Central Limit Theorem (CLT) is a fundamental concept in statistics that describes the distribution of sample means for a sufficiently large sample, regardless of the shape of the original population distribution.

Central Limit Theorem:

For a random sample of size n drawn from any population with a finite mean μ and a finite standard deviation σ , the distribution of the sample means will approach a normal distribution as n becomes sufficiently large. Specifically, as n approaches infinity, the distribution of the sample means will have a mean equal to the population mean ($\mu_{\bar{X}} = \mu$) and a standard deviation equal to the population standard deviation divided by the square root of the sample size ($\sigma_{\bar{X}} = \frac{\sigma}{\sqrt{n}}$).

The Central Limit Theorem is particularly powerful because it allows statisticians to make inferences about population parameters based on the distribution of sample means, **even when the original population distribution is unknown or not normally distributed.**

This theorem forms the basis for many statistical techniques and hypothesis tests that rely on the normal distribution.

	A	B	C	D	E	F	G	H	I	J	K	L
1	y	x				SUMMARY OUTPUT						
2	7.238462	0.025641										
3	6.310256	0.051282				Regression Statistics						
4	8.315385	0.076923				Multiple R	0.906270151					
5	4.787179	0.102564				R Square	0.821325586					
6	5.592308	0.128205				Adjusted R Square	0.819483582					
7	7.830769	0.153846				Standard Error	1.882513522					
8	9.902564	0.179487				Observations	99					
9	5.607692	0.205128				ANOVA						
10	5.146154	0.230769					df	SS	MS	F	Significance F	
11	4.784615	0.25641				Regression	1	1580.159507	1580.16	445.8869	4.72338E-38	
12	7.05641	0.282051				Residual	97	343.7541446	3.543857			
13	9.394872	0.307692				Total	98	1923.913652				
14	6.8	0.333333										
15	4.871795	0.358974										
16	5.376923	0.384615					Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%
17	10.71538	0.410256				'b' → Intercept	5.922586409	0.381284372	15.53325	4.65E-28	5.165842474	6.679330343
18	11.55385	0.435897				'a' → x	5.452241187	0.258203842	21.11603	4.72E-38	4.939778036	5.964704338
19	9.258974	0.461538										
20	8.097436	0.487179										
21	12.10256	0.512821										

$$y = b + a \cdot x$$

How to interpret these values?

What do they mean?

The problem we are trying to solve:

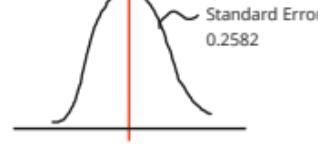
- We have a dataset (x, y) and we know that it is a **random, yet representative sample** (let's call it sample **s1**) of the population
- Based on **s1** we have fitted an SLR model, $y = 5.4522 * x + 5.9226$ by **estimating** the coefficients **a1=5.4522, b1=5.9226** (by minimizing SSE)
- If we had a different random sample **s2**, the calculated values **a2, b2** would have been different. In general, a random sample **si** will result in coefficient values **ai, bi** ...
- So, statistically, the calculated coefficient **ai** can possibly assume different values, depending on the **random sample si** that we get for analysis
- There is an important question that needs to be answered: **What is the probability that the calculated value of ai will be close to or equal to ZERO for some of the si?**
- Why is this a critical question?
 - In the regression model $y = a * x + b$, If **a** is **ZERO** then we do not really have a regression model - ie. y cannot really be predicted in terms of x !
- In the context of our example, the calculated value of **a = 5.4522 ..**
 - What is the probability that this value is obtained **by chance**, due to the peculiarity of sample **s1**?
 - What is the probability that other random samples, **si**, will result in values of **ai** that are very close or equal to **ZERO** - thereby making the model invalid?
 - (*BTW, why only **a**? Why are we not much concerned about **b**?*)
- Can we prove, based only on **s1**, that for any other **si** the calculated value of **ai** has a very high probability of being closer to **5.4522**, and almost never close to **ZERO - thereby establishing the validity of the model?**
 - More practically, can we prove, based only on **s1**, that for any other **si** the calculated value of **ai** has more than 95% probability of being closer to 5.4522, and less than 5% probability of being close to ZERO? (*BTW, why are we talking about 95% and 5%, and not 100% and 0%?*)
 - **So, finally, it comes down to whether we can predict the spread of the values of ai based on just s1 and a1**
 - The **Sampling Distribution of coefficient 'a'** and the **Central Limit Theorem** help us to calculate this probability and, thus, decide about the validity of the **s1** based Regression model.

We will establish the **sampling distribution** of coefficients **a,b** based on s1, a1, b1 as follows:

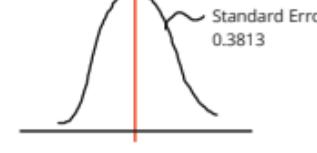
- Based on **s1** calculate the coefficients **a1** (= **5.4522**) and **b1** (= **5.9226**) - both these are **statistics** and it is our goal to check if these values are obtained by chance, or they truly represent the values **obtainable** from most other samples as well.
 - The Sampling Distribution of any **statistic** provides us the mechanism to make this assessment
- To establish the Sampling Distributions of coefficients **a** and **b** we need to find out their standard deviations - the **Standard Errors**. In this context, note the following:
 - Our earlier discussion, we covered the **Sampling Distribution of the Sample Mean**. However, now our object of study is **not** the sample mean but the coefficients **a1, b1** which are calculated from **s1** using optimization methods. Hence the earlier standard error formula ($= \sigma/\sqrt{n}$) does not hold in this case.
- The Standard Error related to coefficients a1 and b1 are given by the formulae:

$$SE(a) = \sqrt{\frac{\hat{\sigma}^2}{\sum_{i=1}^n (x_i - \bar{x})^2}} \quad SE(b) = \sqrt{\frac{\hat{\sigma}^2}{n} \left(1 + \frac{1}{\sum_{i=1}^n (x_i - \bar{x})^2} \right)} \quad \hat{\sigma}^2 \text{ is the estimated variance of the error term in the model.}$$

- On plugging the sample data (**s1**) into these formulae, the Standard Errors of the sampling distribution of **a** and **b** are 0.2582 and 0.3813, respectively, and we know that both Sampling Distributions follow the Normal Distribution.



Sampling Distribution of coefficient a

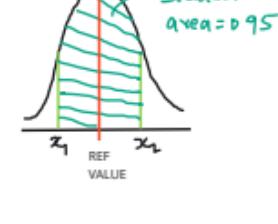


Sampling Distribution of coefficient b

In the formulae alongside we need the value of 'sigma' the population standard deviation. Wherefrom do we get it?

We rely on our sample 's1' for that!

Recollect that s1 is supposed to be representative of the population. Hence, we can assume that the standard deviation of the sample is close to the standard deviation of the population. As we are using this value for estimating the prediction error of a and b, this assumption is acceptable.

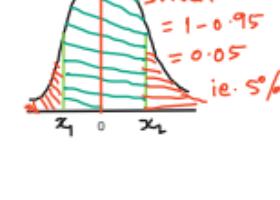


- Lets again recollect that the Sampling Distributions reflect the variations in the **statistics** calculated using different random samples **s1**.
- Lets also recollect that the area under the Sampling Distribution curve (formally known as the **Probability Density Function**) represents probability. For example, the shaded area in the figure alongside gives the probability of the values between x1 and x2 - also note that the **total area** under the probability distribution curve is 1.0
- For any **REF VALUE** (we will use **ZERO** - why?), and the above calculated **STANDARD ERROR** values, we can find out the limits x1 and x2 using Normal Distribution Tables / Python functions / Excel formula, etc., such that the **shaded area under the curve is 0.95**.

- The interval {x1, x2} is said to constitute the **95% Confidence Interval (95% CI)**.
- The interpretation is that all values between x1 and x2 are **statistically not really different from the REFERENCE VALUE**.
- Why? - Because, if the population parameter equals the REFERENCE VALUE, and If you were to take 100 random samples (s1) from the population, 95 of those samples will result in the calculated value of the **statistic** lying between x1 and x2.

- Now we come back to our critical question, posed earlier, and ask it in the form of the following **Statistical Hypothesis**: "**Is a1 statistically different from ZERO**"

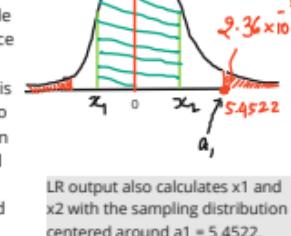
- We want to check if the statistic **a1** is significantly different from **ZERO**
- Hence we choose the reference value to be ZERO,**
- Now, does the 95% CI {x1, x2} include the calculated value of the coefficient a1 (ie. 5.4522)?**
- If YES, then the calculated value 5.4522 is **obtained by chance** and it is **NOT statistically different from ZERO** - hence our regression model is not valid.
- If NO, then we can confidently say the following:
 - That the calculated value 5.4522 is **SIGNIFICANTLY DIFFERENT** from ZERO.
 - That, the value of 5.4522 is not resulting from chance because of the peculiarity of the specific sample **s1**
 - That 95 out of 100 random data samples would also have resulted in a value of a1 closer to 5.4522.
- All this is equivalent to checking whether the calculated value (eg. a1 = 5.4522) lies within the green region (the 95% region) around ZERO or, does it lie in the red region (the 5% region)?
 - If within the 95% (green) region: The calculated value is statistically not different from the reference value ZERO
 - If outside the 95% region (ie. within the remaining 5% region): The calculated value is statistically different from the reference value ZERO - and hence stated to be **SIGNIFICANT, RELEVANT** and **VALID**



- So, finally, **what is p-value and how should it be interpreted?**

- We would like to check what is the position of **a1** (= 5.4522) with respect to the reference value (ZERO) of the established Sampling Distribution of **a**, and this we do by calculating the sum of the area(s) under the curve to the far sides of a1, and then comparing it to the reference value of 0.05 (ie. 5%).
- p-value** is defined as the **sum** of the areas under the curve beyond the far side of the calculated value, and the corresponding area on the opposite side - since a1 can take positive and negative values. For example, in the figure alongside, the red shaded area to the right lies beyond the calculated value of a1, and this area is 2.36E-38. Similar sized area is also considered to the left, and these two add to a total of 4.72E-38, which is the p-value value reported in the regression output. This belongs to the class of Hypothesis Tests known as the **two-tailed** tests because a1 can take both positive and negative values.

- In some Hypothesis Tests, the region of interest is on only one side, and such tests are known as **one-tailed** tests.
- As this p-value is much less than the threshold value of 0.05 (ie. 5%), we can conclude that 5.4522 is outside the 95% confidence interval for the reference value of ZERO, hence it **STATISTICALLY SIGNIFICANT** (ie. statistically different from ZERO) and hence valid.
- If the p-value was greater than 0.05, it would necessarily imply that the calculated point 5.4522 is actually lying inside the 95% CI (the green zone), hence **NOT STATISTICALLY SIGNIFICANT** (ie. not different from ZERO).
- As we will see later, in ML, the p-values help us to **select** the most appropriate independent variables (**x's**) to be included in the model.



LR output also calculates x1 and x2 with the sampling distribution centered around a1 = 5.4522. In the example:

x1 = 4.9398

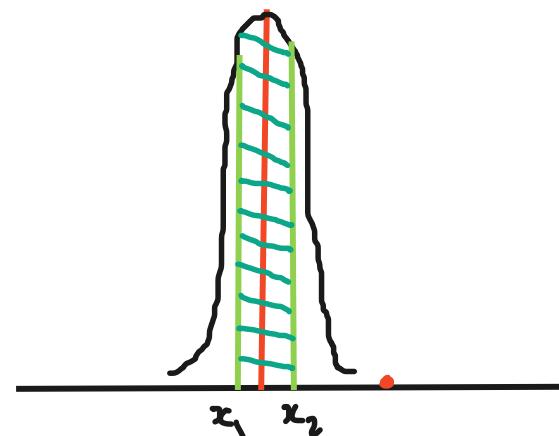
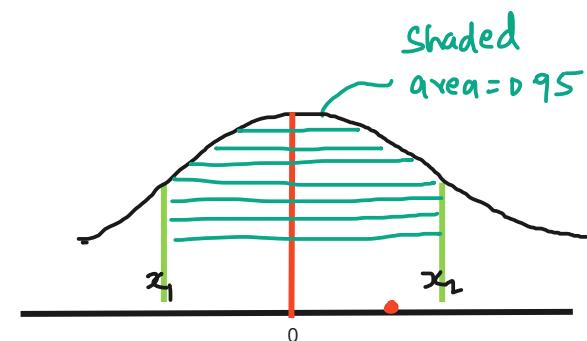
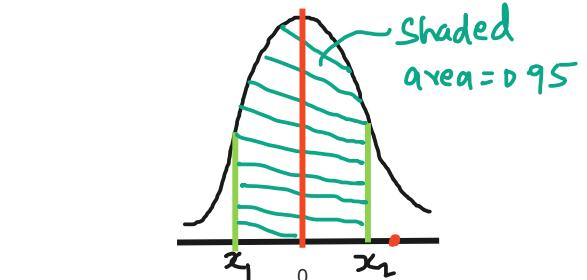
x2 = 5.9647

It is worth noting that ZERO is NOT a part of this interval.

This also tells us that 95 out of 100 samples s1's results can fit within this interval.

The implications of sample size on the confidence interval:

- It should be noted that the formulae for standard errors have the sample size **n** in the denominator.
- Consequently, **if the sample size is reduced** the standard error value becomes larger and the sampling distribution curve gets stouter and wider as shown alongside, and the x_1 and x_2 values get pushed further away from ZERO
- The net effect is that values such as the red dot, which was SIGNIFICANT earlier now falls within the **not statistically different from ZERO zone!** The zone of uncertainty has widened, and it is more difficult to get a valid model.
- Conversely, if the sample size is increased, the standard error value becomes smaller and the sampling distribution curve gets slimmer and the x_1 , x_2 values get pulled towards ZERO - the range of values considered *statistically equal to ZERO* becomes small.
 - This increases the chance that the red dot (calculated a_1) will lie far outside the 95% CI for ZERO and hence it will be considered as statistically significant. In effect, **uncertainty in the model reduces as the sample size increases.**



What about **F** and *Significance F*?

- These are relevant in case of Multiple Linear Regression (MLR) where $y = f(x_1, x_2, x_3, x_4, \dots)$
- The **F-statistic** is defined as follows:

$$F = \frac{\frac{SSR}{n-k-1}}{\frac{SSE}{n-1}}$$

n is the number of observations,

k is the number of predictors (independent variables) in the model.

- As we can see, it is the ratio of **average variance explained by regression** and **average variance attributable to random errors (MSR / MSE)**
- The better the regression model, the larger will be the value of the **F-statistic** - that is, significantly more variance in the data will be explained by the regression model.
- **Significance F** is nothing but the p-value associated with the **F-statistic**. Therefore, if **Significance F** is less than 0.05 (5%) we can rely on the calculated **F-statistic** and consider it **statistically significant** and use it to confidently make judgements about the overall quality of the Linear Regression model.
- **We assess the LR model quality by taking into account the calculated coefficients (and their p-values) and the F-statistic (and its p-value)**
- This is not the end !! There are a few more metrics, that we will encounter soon ...

ANOVA						
	df	ss	MS	F	Significance F	
Regressor	4	85.75865	21.43966	222.0403	1.26E-67	
Residual	176	16.99412	0.096558			
Total	180	102.7528				

	Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%
Intercept	0.31406	0.11176	2.810135	0.005513	0.093498	0.534622
x1	12.33273	1.5577	7.917267	2.62E-13	9.258555	15.40691
x2	-38.302	6.359842	-6.02247	9.77E-09	-50.8533	-25.7506
x3	30.31208	9.568272	3.167978	0.00181	11.42877	49.19539
x4	-4.00187	4.746218	-0.84317	0.400277	-13.3687	5.36495

OLS Regression Results (ORDINARY LEAST SQ.)

=====
 Dep. Variable: y R-squared: 0.891
 Model: OLS Adj. R-squared: 0.890
 Method: Least Squares F-statistic: 607.6
 Date: Tue, 23 Jan 2024 Prob (F-statistic): 2.04e-37
 Time: 22:20:19 Log-Likelihood: -102.30
 No. Observations: 76 AIC: 208.6
 Df Residuals: 74 BIC: 213.3
 Df Model: 1
 Covariance Type: nonrobust
 =====

	coef	std err	t	P> t	[0.025	0.975]
const	3.2028	0.221	14.499	0.000	2.763	3.643
x	9.1104	0.370	24.650	0.000	8.374	9.847

 =====

Omnibus:	5.816	Durbin-Watson:	1.896
Prob(Omnibus):	0.055	Jarque-Bera (JB):	2.579
Skew:	-0.112	Prob(JB):	0.275
Kurtosis:	2.126	Cond. No.	4.42

 =====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OUTPUT CREATED BY
 OLS function of
 STATSMODELS library of
 Python .

We have already encountered some of the generated numbers like R2, F-statistic, etc. But the adjacent block of output contains many others ... what are they, why are they important, and how to interpret them?

AIC (Akaike Information Criteria) and BIC (Bayes Information ...)

- Both these are estimators of **prediction error**. They help in model selection.
- These numbers are used to compare across models. **A lower number indicates a better model**
- A difference in AIC or BIC value of 2, between models being compared, is considered significant. The model with a lower AIC or BIC value is designated as the better model, and becomes a candidate for selection.

- Omnibus statistic:** This is a numeric value calculated from the skewness and kurtosis of the residuals (the difference between the predicted and actual values). A low value suggests the residuals are closer to a normal distribution, while a high value indicates deviation from normality.

- Omnibus p-value:** This value represents the probability of observing the calculated Omnibus statistic, assuming the null hypothesis of normally distributed residuals is true. A low p-value (typically < 0.05) suggests there is significant evidence to reject the null hypothesis, implying the residuals are not normally distributed.

Skewness measures the asymmetry of the probability distribution of a dataset. For a normal distribution, the skewness should be **close to 0**.

- Skewness = 0:** The data is symmetric, which is a characteristic of a normal distribution.
- Skewness > 0:** The data is positively skewed (right-tailed), meaning the tail on the right side is longer or fatter.
- Skewness < 0:** The data is negatively skewed (left-tailed), meaning the tail on the left side is longer or fatter.

Interpretation:

- If the skewness is significantly different from 0, the data is unlikely to follow a normal distribution.

Kurtosis measures the “tailedness” of the probability distribution, indicating whether the data has heavy or light tails compared to a normal distribution. For a normal distribution, the kurtosis is **close to 3** (or 0 if using excess kurtosis).

- Kurtosis = 3 (or Excess Kurtosis = 0):** The data has tails similar to a normal distribution.
- Kurtosis > 3 (or Excess Kurtosis > 0):** The data has heavier tails (more outliers) than a normal distribution (leptokurtic).
- Kurtosis < 3 (or Excess Kurtosis < 0):** The data has lighter tails (fewer outliers) than a normal distribution (platykurtic).

Interpretation:

- If the kurtosis is significantly different from 3 (or excess kurtosis significantly different from 0), the data is unlikely to follow a normal distribution.

This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

Jarque-Bera Test

- In the context of Ordinary Least Squares (OLS) regression, the Jarque-Bera test is used to **check the normality of the residuals.**
 - **Residuals**, the difference between the predicted and actual values in your model, play a crucial role in OLS analysis.
 - Their normality is one of the key assumptions for the validity of statistical inferences drawn from the model.
- A **low statistic** : Low value of the Jarque-Bera statistic (< 2) along with high p-value (ie. > 0.05) indicate that the residuals follow Normal Distribution.
- A **high statistic** : High value of the Jarque-Bera statistic (> 6) often accompanied with low p-values (ie. < 0.05) indicate that the residuals DO NOT follow Normal Distribution.

Durbin-Watson Test

- In the context of Ordinary Least Squares (OLS) regression, the Durbin-Watson (DW) test is a diagnostic tool used to check for **autocorrelation** in the residuals (errors) of the model.

Autocorrelation occurs when there's a dependence between subsequent errors, meaning the error term at one point in time influences the error term at another point.

- Its value always falls between 0 and 4, with specific interpretations:
 - **2.0:** Indicates no autocorrelation (ideal scenario).
 - **0 to less than 2.0:** Suggests positive autocorrelation (errors tend to cluster together, either positive or negative).
 - **More than 2.0 to 4:** Suggests negative autocorrelation (errors tend to alternate between positive and negative).
- This test is used as a first check, and not a definitive test.

Good regression!

Condition Number

The Condition Number in Ordinary Least Squares (OLS) refers to a **measure of how sensitive the estimated coefficients are to small changes in the data**. It's not directly related to any specific variable or error term, but rather evaluates the overall stability and robustness of the model's solution. Its calculation is based on eigenvalues

- A **low Condition Number** indicates that the coefficients react minimally to small changes in the data (stable, robust model).
- A **high Condition Number** signifies that even slight data variations can significantly alter the coefficients (sensitive, potentially unstable model).

Why is it important?

- A high Condition Number suggests the model might be fitting noise or capturing spurious relationships due to its sensitivity to slight data changes. This makes the estimated coefficients less reliable and conclusions less trustworthy.
- In extreme cases, a very high Condition Number can lead to numerical issues during calculations, rendering the model estimation altogether unstable.

Interpretation:

There's no single threshold for a "good" or "bad" Condition Number. However, in general:

- **Values below 10** are considered acceptable, indicating a relatively stable model.
- **Values above 30** raise concerns about sensitivity and potential instability.
- **Values above 100** are a strong indicator of an unreliable model requiring further investigation or improvement.

Output	Interpretation	Specific Limits/Values (if applicable)
R^2	Proportion of variance in the dependent variable explained by the model.	Range: [0, 1], higher values are desirable.
Adjusted R^2	R^2 adjusted for the number of predictors; a measure of model fit.	Like R^2, but adjusted for model complexity.
F-statistic	Tests the overall significance of the regression model.	Critical values based on significance level (e.g., 0.05).
AIC (Akaike's IC)	A measure of model goodness-of-fit, balancing complexity and fit.	Lower values are better; used for model comparison.
BIC (Bayesian IC)	Similar to AIC but penalizes model complexity more heavily.	Lower values are better; stricter penalty for complexity.
Log Likelihood	A measure of how well the model explains the observed data.	Higher values indicate better model fit.
Omnibus	Refers to a specific statistic and its associated p-value that test the normality of the residuals. It is a combination of multiple tests like Jarque-Bera test, Shapiro-Wilkes test and Kolmogorov-Smirnov test.	For the residuals to have Normal Distribution, the Omnibus statistic should have low value and p-value should be > 0.05
Durbin-Watson test	Tests for autocorrelation in the residuals; values around 2 suggest no autocorrelation.	Range: [0, 4], close to 2 indicates no significant autocorrelation.
Jarque-Bera test	Tests for normality of residuals based on skewness and kurtosis. Test statistic value closer to zero implies residuals are normally distributed NULL Hypothesis: The residuals are Normally Distributed	Critical values based on significance level (e.g., 0.05). If p-value < 0.05, then NULL Hypothesis is rejected implying the residuals are NOT normally distributed. If residuals are normally distributed, p-value > 0.05.
Condition Number	Measures sensitivity to changes in input variables; high values indicate multicollinearity.	No strict limits; values above 30 may indicate multicollinearity.
Skew	A measure of the asymmetry of the residuals distribution.	Range: $(-\infty, \infty)$; 0 for a perfectly symmetric distribution.
Kurtosis	A measure of the "This file is meant for personal use by mepravintpatil@gmail.com only." Sharing or publishing the contents in part or full is liable for legal action.	Range: $(-\infty, \infty)$; 3 for a normal distribution (excess kurtosis).

MULTIPLE LINEAR REGRESSION.

- SLR deals with only one independent variable, and takes the form:

$$y = a \cdot x + b \quad \text{or} \quad y = \beta_0 + \beta_1 x$$

- MLR deals with more than one independent variable, and takes the form:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k$$

here $[x_1, x_2, x_3, \dots, x_k]$ are the independent variables, also known as "features"

A data set for MLR will look as shown below

y	x1	x2	x3	x4
0.038117	0	0	0	0
0.896468	0.005556	3.09E-05	1.71E-07	9.53E-10
0.159546	0.011111	0.000123	1.37E-06	1.52E-08
0.863764	0.016667	0.000278	4.63E-06	7.72E-08
1.106349	0.022222	0.000494	1.1E-05	2.44E-07
1.010169	0.027778	0.000772	2.14E-05	5.95E-07
0.278498	0.033333	0.001111	3.7E-05	1.23E-06
1.114231	0.038889	0.001512	5.88E-05	2.29E-06
1.029804	0.044444	0.001975	8.78E-05	3.9E-06
0.37387	0.05	0.0025	0.000125	6.25E-06
0.971634	0.055556	0.003086	0.000171	9.53E-06
0.975377	0.061111	0.003735	0.000228	1.39E-05
1.079774	0.066667	0.004444	0.000296	1.98E-05
1.24279	0.072222	0.005216	0.000377	2.72E-05
0.644699	0.077778	0.006049	0.000471	3.66E-05
0.656177	0.083333	0.006944	0.000579	4.82E-05

features.

TRAIN DATA.

This file is meant for personal use by megovintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

In MLR, the goal is to express 'y' as a linear combination of x_1, x_2, \dots

$$\therefore y_1 = \beta_0 + \beta_1 x_{11} + \beta_2 x_{21} + \dots + \beta_k x_{k1} + e_1$$

$$y_2 = \beta_0 + \beta_1 x_{12} + \beta_2 x_{22} + \dots + \beta_k x_{k2} + e_2$$

$$\vdots$$
$$y_m = \beta_0 + \beta_1 x_{1m} + \dots + \beta_k x_{km} + e_m$$

'k' features

What is MLR?

Using the values of all x_{ij} and the corresponding values of y_j , find out the most appropriate $\beta_0, \beta_1, \dots, \beta_k$.

How do we go about it?

The model that we create, ie. the values of β_0, β_1 , etc. that we identify should be such that

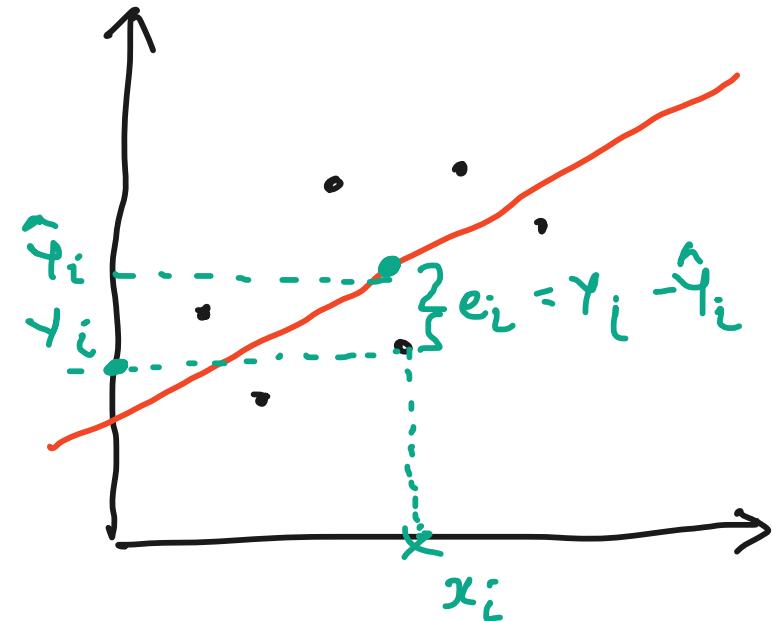
$$\sum_{i=1}^n e_i^2$$

should be minimized
(minimize the sum of square of errors, SSE)

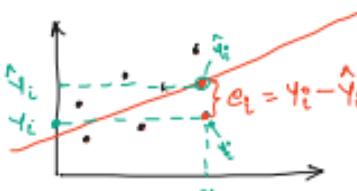
Note:

$$y_i = \beta_0 + \sum_{j=1}^m \beta_j x_{ij} + e_i$$

$$\hat{y}_i = \beta_0 + \sum_{j=1}^m \hat{\beta}_j x_{ij}$$



MLR GRADIENT DESCENT (OPTIONAL READING!)



$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \dots + \beta_k x_{ik}$$

Matrices used in the derivations...

$$\begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}_{m \times 1} = \begin{bmatrix} 1 & x_{11} & x_{12} & x_{13} & \dots & x_{1k} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m1} & x_{m2} & x_{m3} & \dots & x_{mk} \end{bmatrix}_{m \times k} \begin{bmatrix} \beta_0 \\ \vdots \\ \beta_k \end{bmatrix}_{k \times 1} + \begin{bmatrix} e_1 \\ \vdots \\ e_m \end{bmatrix}_{m \times 1}$$

m = Number of observations (records)

k = Number of features (independent variables)

$l = k+1$

$$\rightarrow y = X \cdot \beta + E \quad \text{and} \quad \hat{y} = X \cdot \beta \quad \text{... the regression 'line'}$$

$$E = y - \hat{y}$$

$$E^T \cdot E = (y - \hat{y})^T (y - \hat{y}) \quad \text{... Sum of Squares of Errors}$$

$$J = \frac{1}{2m} (E^T \cdot E) \quad \text{... Cost function.}$$

Goal: minimize J (ie: make $\frac{\partial J}{\partial \beta} = 0$)

$$J = \frac{1}{2m} (y - \hat{y})^T (y - \hat{y})$$

$$= \frac{1}{2m} [(y^T - \beta^T x^T)(y - \hat{y})]$$

$$= \frac{1}{2m} [y^T \cdot y - y^T (x \beta) - (\beta^T x^T) \cdot y + (\beta^T x^T) \cdot (x \beta)]$$

$$\frac{\partial J}{\partial \beta} = \frac{1}{2m} \left[\frac{d}{d \beta} \{y^T \cdot y\} - \frac{d}{d \beta} \{y^T (x \beta)\} - \frac{d}{d \beta} \{(\beta^T x^T) \cdot y\} + \frac{d}{d \beta} \{(\beta^T x^T) \cdot (x \beta)\} \right]$$

$$A = \frac{d}{d \beta} \{y^T \cdot y\} = 0 \quad \text{... } y \text{ is a vector of constants.}$$

$$B = \frac{d}{d \beta} \{y^T (x \beta)\} = y^T \cdot \frac{d}{d \beta} (x \beta) + \frac{d}{d \beta} y^T \cdot (x \beta)$$

$$= y^T x + 0 = y^T x \quad \text{... Row vector } (1 \times l)$$

$$C = \frac{d}{d \beta} \{(\beta^T x^T) \cdot y\} = \beta^T x^T \frac{d}{d \beta} y + \frac{d}{d \beta} (\beta^T x^T) \cdot y$$

$$= 0 + x^T y \quad \text{... Column Vector } (l \times 1)$$

$$D = \frac{d}{d \beta} \{(\beta^T x^T) \cdot (x \beta)\} = (\beta^T x^T) \frac{d}{d \beta} (x \beta) + \frac{d}{d \beta} (\beta^T x^T) \cdot (x \beta)$$

$$= (\beta^T x^T x)_{l \times l} + (x^T x \beta)_{l \times 1}$$

$$\frac{\partial J}{\partial \beta} = \frac{1}{2m} \left[(\beta^T x^T x) + (x^T x \beta) - y^T x - x^T y \right]$$

There are all 'vectors' and the results in the pairs shown below are equal in values. Hence re-arranging and simplifying ...

$$= \frac{1}{2m} \left[2 \cdot x^T x \beta - 2 x^T y \right]_{(l \times 1)}$$

$$= x^T (x \beta - y) \quad \dots$$

... Since $\hat{y} = X \beta$

$$\frac{\partial J}{\partial \beta} = x^T (\hat{y} - y)$$

The Gradient Descent process

① assume some value for β . Eg: $\begin{bmatrix} \beta_0 \\ \vdots \\ \beta_k \end{bmatrix} \leftarrow$

② Using $\hat{y} = X \beta$, evaluate $\hat{y} \leftarrow$

③ Calculate $\frac{\partial J}{\partial \beta}$ as per the above expression

④ by assuming some value for η (eg: 0.05)

⑤ calculate new values for β using the following expression

$$\beta_{new} \leftarrow \beta_{old} - \eta (\nabla J)$$

⑥ repeat steps 2 to 5 till $\|\beta_{new} - \beta_{old}\| < threshold$ (and $\eta > 0.01$)



MULTIPLE LINEAR REGRESSION.(MLR)

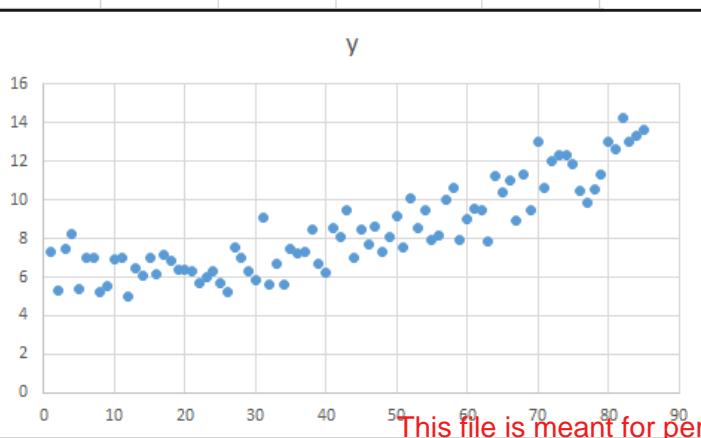
- In MLR, more than one independent variable x_i potentially determine the dependent variable 'y'

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \dots + \beta_K x_K$$

- MLR involves calculating the coefficients β_i using the given dataset (TRAIN DATA).
- As in the case of SLR, these values are obtained by minimizing the SSE (explained in a separate document)
- MLR involves identifying the most relevant predictors as explained subsequently.

MLR using the dataset **data-set-for-MLR.xlsx**

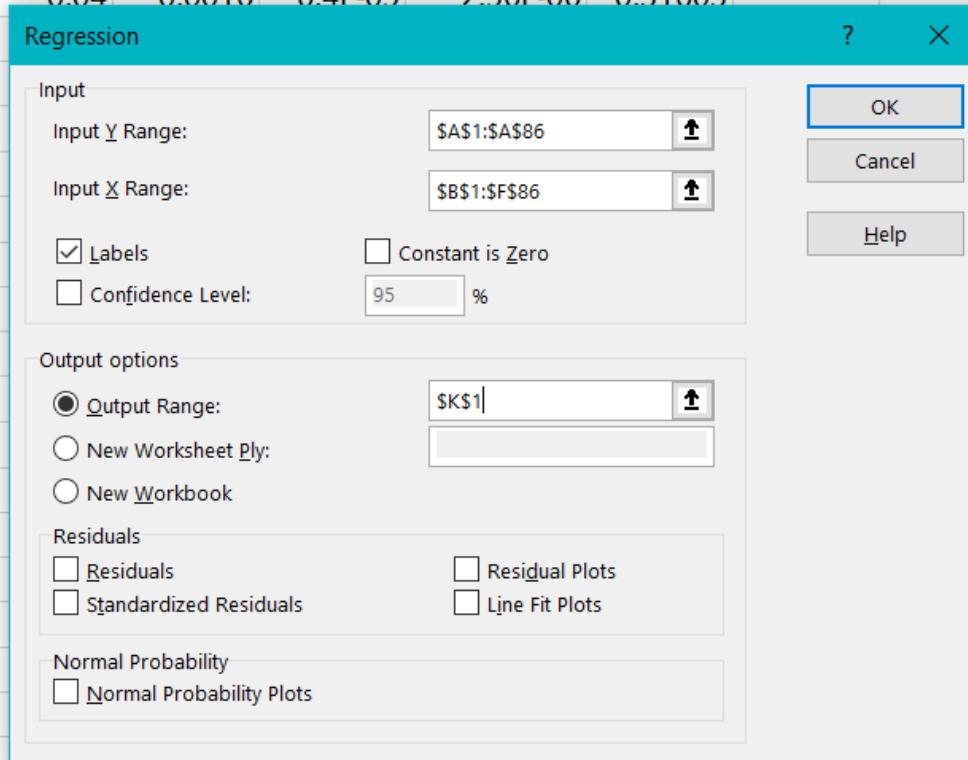
	A	B	C	D	E	F
1	y	x1	x2	x3	x4	x5
2	7.29594	0	0	0	0	0.56109
3	5.30545	0.02	0.0004	8E-06	1.6E-07	0.89668
4	7.42688	0.03	0.0009	2.7E-05	8.1E-07	0.9675
5	8.2255	0.04	0.0016	6.4E-05	2.56E-06	0.31603
6	5.3746	0.05	0.0025	0.00013	6.25E-06	0.74414
7	7.02144	0.06	0.0036	0.00022	1.296E-05	0.19197
8	6.98843	0.07	0.0049	0.00034	2.401E-05	0.86862
9	5.21817	0.08	0.0064	0.00051	4.096E-05	0.74443
10	5.55326	0.09	0.0081	0.00073	6.561E-05	0.801
11	6.94546	0.1	0.01	0.001	0.0001	0.05507
12	7.02019	0.11	0.0121	0.00133	0.0001464	0.61864
13	5.03213	0.12	0.0144	0.00173	0.0002074	0.80112
14	6.49254	0.13	0.0169	0.0022	0.0002856	0.78146
15	6.0491					
16	6.99021					
17	6.13994					
18	7.16849					
19	6.84618					
20	6.4126					
21	6.40522					
22	6.34211					
23	5.71266					
24	5.99409					
25	6.29404					



- The dataset in the file consists of the **train** dataset of 85 observations and the **test** dataset of 15 observations
- We will create an MLR model using the train dataset and subsequently validate the model using the test dataset
- We start by creating an MLR model using all the **x** variables (also known as **features**)
- A scatter plot of **y** reveals that the observations are non-linear ...
 - So, will **Linear Regression** be able to create a good and acceptable model??

	A	B	C	D	E	F	I	J
1	y	x1	x2	x3	x4	x5		
2	7.29594		0	0	0	0	0.56109	
3	5.30545		0.02	0.0004	8E-06	1.6E-07	0.89668	
4	7.42688		0.03	0.0009	2.7E-05	8.1E-07	0.9675	
5	8.2255		0.04	0.0016	6.4F-05	2.56F-06	0.31603	
6	5.3746							
7	7.02144							
8	6.98843							
9	5.21817							
10	5.55326							
11	6.94546							
12	7.02019							
13	5.03213							
14	6.49254							
15	6.0491							
16	6.99021							
17	6.13994							
18	7.16849							
19	6.84618							
20	6.4126							
21	6.40522							
22	6.34211							
23	5.71266	0.24	0.0576	0.01382	0.0033178	0.29455		
24	5.99409	0.25	0.0625	0.01563	0.0039003	0.85295		
25	6.29404	0.26	0.0676	0.01758	0.0045698	0.0629		

Invoking the Linear Regression functionality of Excel and selecting the variables ...



This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

	A	B	C	D	E	F	I	J	K	L	M	N	O	P	Q
1	y	x1	x2	x3	x4	x5			SUMMARY OUTPUT						
2	7.29594	0	0	0	0	0.56109									
3	5.30545	0.02	0.0004	8E-06	1.6E-07	0.89668			Regression Statistics						
4	7.42688	0.03	0.0009	2.7E-05	8.1E-07	0.9675			Multiple R	0.912936908					
5	8.2255	0.04	0.0016	6.4E-05	2.56E-06	0.31603			R Square	0.833453799	- OK				
6	5.3746	0.05	0.0025	0.00013	6.25E-06	0.74414			Adjusted R Square	0.8229129					
7	7.02144	0.06	0.0036	0.00022	1.296E-05	0.19197			Standard Error	0.991752189					
8	6.98843	0.07	0.0049	0.00034	2.401E-05	0.86862			Observations	85					
9	5.21817	0.08	0.0064	0.00051	4.096E-05	0.74443			ANOVA						
10	5.55326	0.09	0.0081	0.00073	6.561E-05	0.801				df	SS	MS	F	Significance F	
11	6.94546	0.1	0.01	0.001	0.0001	0.05507			Regression	5	388.848	77.7697	79.0686	2.7E-29	→ OK
12	7.02019	0.11	0.0121	0.00133	0.0001464	0.61864			Residual	79	77.7022	0.98357			
13	5.03213	0.12	0.0144	0.00173	0.0002074	0.80112			Total	84	466.551				
14	6.49254	0.13	0.0169	0.0022	0.0002856	0.78146									
15	6.0491	0.15	0.0225	0.00338	0.0005063	0.68791									
16	6.99021	0.16	0.0256	0.0041	0.0006554	0.59236									
17	6.13994	0.17	0.0289	0.00491	0.0008352	0.20935									
18	7.16849	0.18	0.0324	0.00583	0.0010498	0.96488			Intercept	6.908046001	0.59691	11.5731	1.1E-18	5.71993	8.09616
19	6.84618	0.19	0.0361	0.00686	0.0013032	0.48896			x1	-9.846555611	7.41707	-1.32755	0.18815	-24.6099	4.91676
20	6.4126	0.21	0.0441	0.00926	0.0019448	0.50876			x2	39.10941696	30.7379	1.27235	0.20698	22.0728	100.292
21	6.40522	0.22	0.0484	0.01065	0.0023426	0.71971			x3	-42.57714279	46.9855	-0.90618	0.3676	136.099	50.9451
22	6.34211	0.23	0.0529	0.01217	0.0027984	0.49221			x4	20.26808239	23.6657	0.85643	0.39435	-26.8374	67.3736
23	5.71266	0.24	0.0576	0.01382	0.0033178	0.29455			x5	0.10562362	0.41144	0.25672	0.79806	-0.71333	0.92457
24	5.99409	0.25	0.0625	0.01563	0.0039063	0.85295									
25	6.29404	0.26	0.0676	0.01758	0.0045698	0.0629									

None of these values are acceptable
since their corresponding p-values
are all MUCH GREATER than 0.05

So, we DISCARD x_5 - which has
the highest p-value, and re-create
the model.

	A	B	C	D	E	F	I	J	K	L	M	N	O	P	Q
1	y	x1	x2	x3	x4	x5			SUMMARY OUTPUT						
2	7.29594	0	0	0	0	0.56109									
3	5.30545	0.02	0.0004	8E-06	1.6E-07	0.89668			Regression Statistics						
4	7.42688	0.03	0.0009	2.7E-05	8.1E-07	0.9675			Multiple R	0.912860812					
5	8.2255	0.04	0.0016	6.4E-05	2.56E-06	0.31603			R Square	0.833314862					
6	5.3746	0.05	0.0025	0.00013	6.25E-06	0.74414			Adjusted R Square	0.824980605					
7	7.02144	0.06	0.0036	0.00022	1.296E-05	0.19197			Standard Error	0.985945239					
8	6.98843	0.07	0.0049	0.00034	2.401E-05	0.86862			Observations	85					
9	5.21817	0.08	0.0064	0.00051	4.096E-05	0.74443									
10	5.55326	0.09	0.0081	0.00073	6.561E-05	0.801			ANOVA						
11	6.94546	0.1	0.01	0.001	0.0001	0.05507				df	SS	MS	F	Significance F	
12	7.02019	0.11	0.0121	0.00133	0.0001464	0.61864			Regression	4	388.783	97.1959	99.9867	2.6E-30	
13	5.03213	0.12	0.0144	0.00173	0.0002074	0.80112			Residual	80	77.767	0.97209			
14	6.49254	0.13	0.0169	0.0022	0.0002856	0.78146			Total	84	466.551				
15	6.0491	0.15	0.0225	0.00338	0.0005063	0.68791									
16	6.99021	0.16	0.0256	0.0041	0.0006554	0.59236				Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%
17	6.13994	0.17	0.0289	0.00491	0.0008352	0.20935			Intercept	6.982707415	0.51821	13.4746	2.8E-22	5.95143	8.01398
18	7.16849	0.18	0.0324	0.00583	0.0010498	0.96488			x1	-9.95691922	7.36125	-1.35261	0.17999	-24.6063	4.69243
19	6.84618	0.19	0.0361	0.00686	0.0013032	0.48896			x2	39.18955261	30.5563	1.28254	0.20336	-21.6194	99.9986
20	6.4126	0.21	0.0441	0.00926	0.0019448	0.50876			x3	-42.50561677	46.7095	-0.91	0.36556	-135.461	50.4493
21	6.40522	0.22	0.0484	0.01065	0.0023426	0.71971			x4	20.19742663	23.5256	0.85853	0.39316	-26.62	67.0148
22	6.34211	0.23	0.0529	0.01217	0.0027984	0.49221									
23	5.71266	0.24	0.0576	0.01382	0.0033178	0.29455									
24	5.99409	0.25	0.0625	0.01563	0.0039063	0.85295									
25	6.29404	0.26	0.0676	0.01758	0.0045698	0.0629									

discard x_4 and proceed ...

	A	B	C	D	E	F	I	J	K	L	M	N	O	P	Q
1	y	x1	x2	x3	x4	x5			SUMMARY OUTPUT						
2	7.29594	0	0	0	0	0.56109									
3	5.30545	0.02	0.0004	8E-06	1.6E-07	0.89668			Regression Statistics						
4	7.42688	0.03	0.0009	2.7E-05	8.1E-07	0.9675			Multiple R	0.912019254					
5	8.2255	0.04	0.0016	6.4E-05	2.56E-06	0.31603			R Square	0.83177912					
6	5.3746	0.05	0.0025	0.00013	6.25E-06	0.74414			Adjusted R Square	0.825548717					
7	7.02144	0.06	0.0036	0.00022	1.296E-05	0.19197			Standard Error	0.984343752					
8	6.98843	0.07	0.0049	0.00034	2.401E-05	0.86862			Observations	85					
9	5.21817	0.08	0.0064	0.00051	4.096E-05	0.74443			ANOVA						
10	5.55326	0.09	0.0081	0.00073	6.561E-05	0.801				df	SS	MS	F	Significance F	
11	6.94546	0.1	0.01	0.001	0.0001	0.05507			Regression	3	388.067	129.356	133.503	2.9E-31	
12	7.02019	0.11	0.0121	0.00133	0.0001464	0.61864			Residual	81	78.4835	0.96893			
13	5.03213	0.12	0.0144	0.00173	0.0002074	0.80112			Total	84	466.551				
14	6.49254	0.13	0.0169	0.0022	0.0002856	0.78146				Coefficients	Standard Err	t Stat	P-value	Lower 95%	Upper 95%
15	6.0491	0.15	0.0225	0.00338	0.0005063	0.68791			Intercept	6.717731373	0.4156	16.164	4.5E-27	5.89082	7.54464
16	6.99021	0.16	0.0256	0.0041	0.0006554	0.59236			x1	-4.499675935	3.70651	-1.21399	0.22828	-11.8745	2.87513
17	6.13994	0.17	0.0289	0.00491	0.0008352	0.20935			x2	14.05360633	8.73189	1.60946	0.11141	-3.32013	31.4273
18	7.16849	0.18	0.0324	0.00583	0.0010498	0.96488			x3	-2.717152907	5.81602	-0.46718	0.64162	-14.2892	8.8549
19	6.84618	0.19	0.0361	0.00686	0.0013032	0.48896									
20	6.4126	0.21	0.0441	0.00926	0.0019448	0.50876									
21	6.40522	0.22	0.0484	0.01065	0.0023426	0.71971									
22	6.34211	0.23	0.0529	0.01217	0.0027984	0.49221									
23	5.71266	0.24	0.0576	0.01382	0.0033178	0.29455									
24	5.99409	0.25	0.0625	0.01563	0.0039063	0.85295									
25	6.29404	0.26	0.0676	0.01758	0.0045698	0.0629									

Discard x_3 and proceed ...

	A	B	C	D	E	F	I	J	K	L	M	N	O	P	Q
1	y	x1	x2	x3	x4	x5			SUMMARY OUTPUT						
2	7.29594	0	0	0	0	0.56109									
3	5.30545	0.02	0.0004	8E-06	1.6E-07	0.89668			Regression Statistics						
4	7.42688	0.03	0.0009	2.7E-05	8.1E-07	0.9675			Multiple R	0.911770714					
5	8.2255	0.04	0.0016	6.4E-05	2.56E-06	0.31603			R Square	0.831325834					
6	5.3746	0.05	0.0025	0.00013	6.25E-06	0.74414			Adjusted R Square	0.82721183					
7	7.02144	0.06	0.0036	0.00022	1.296E-05	0.19197			Standard Error	0.979640446					
8	6.98843	0.07	0.0049	0.00034	2.401E-05	0.86862			Observations	85					
9	5.21817	0.08	0.0064	0.00051	4.096E-05	0.74443			ANOVA						
10	5.55326	0.09	0.0081	0.00073	6.561E-05	0.801				df	SS	MS	F	Significance F	
11	6.94546	0.1	0.01	0.001	0.0001	0.05507			Regression	2	387.856	193.928	202.072	2E-32	
12	7.02019	0.11	0.0121	0.00133	0.0001464	0.61864			Residual	82	78.695	0.9597			
13	5.03213	0.12	0.0144	0.00173	0.0002074	0.80112			Total	84	466.551				
14	6.49254	0.13	0.0169	0.0022	0.0002856	0.78146									
15	6.0491	0.15	0.0225	0.00338	0.0005063	0.68791									
16	6.99021	0.16	0.0256	0.0041	0.0006554	0.59236									
17	6.13994	0.17	0.0289	0.00491	0.0008352	0.20935									
18	7.16849	0.18	0.0324	0.00583	0.0010498	0.96488									
19	6.84618	0.19	0.0361	0.00686	0.0013032	0.48896									
20	6.4126	0.21	0.0441	0.00926	0.0019448	0.50876									
21	6.40522	0.22	0.0484	0.01065	0.0023426	0.71971									
22	6.34211	0.23	0.0529	0.01217	0.0027984	0.49221									
23	5.71266	0.24	0.0576	0.01382	0.0033178	0.29455									
24	5.99409	0.25	0.0625	0.01563	0.0039063	0.85295									
25	6.29404	0.26	0.0676	0.01758	0.0045698	0.0629									

The model seems ok now ...
 p-value of x_1 is very close
 to 0.05, so we choose to
 keep it ... -

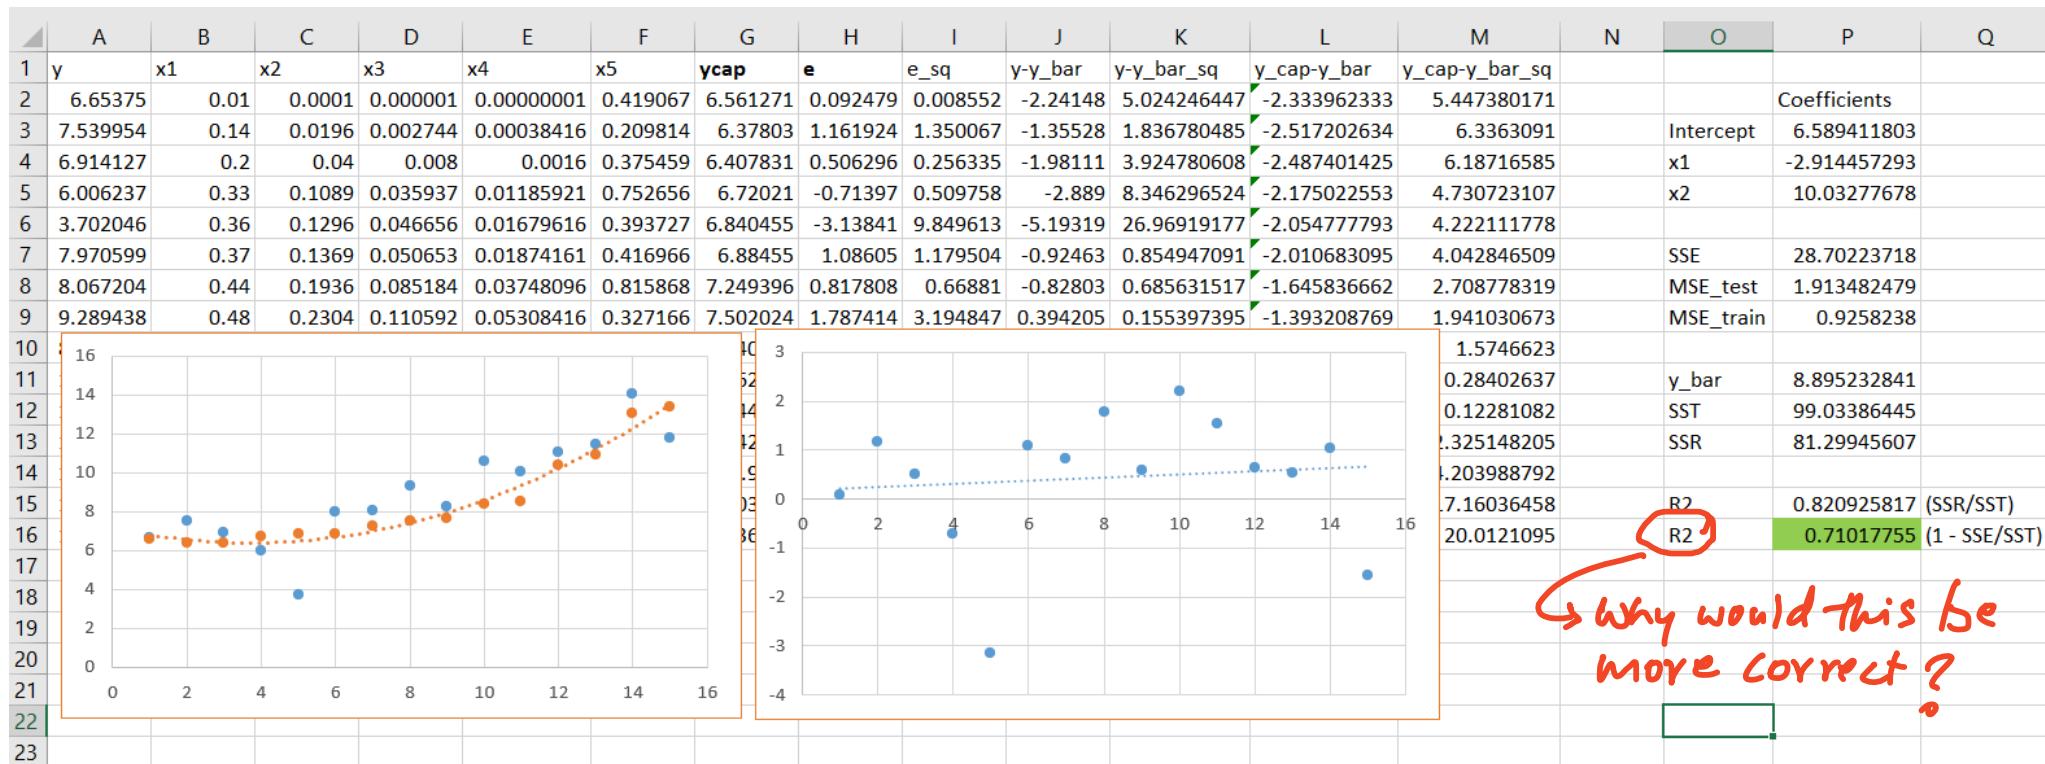
	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	
1	y	x1	x2	x3	x4	x5	ycap	e			SUMMARY OUTPUT							
2	7.295939		0	0	0	0	0.561088	6.589412	0.706527									
3	5.305447	0.02	0.0004	0.000008	0.00000016	0.896684	6.535136	-1.22969			Regression Statistics							
4	7.4	16					5.511008	0.915876			Multiple R	0.911770714						
5	8.2	14					5.488886	1.736618			R Square	0.831325834						
6	5.3	12					5.468771	-1.09418			Adjusted R Square	0.82721183						
7	7.0	10					5.450662	0.570781			Standard Error	0.979640446						
8	6.	8					6.43456	0.55387			Observations	85						
9	5.2	6					5.420465	-1.2023										
10	5.5	4					5.408376	-0.85512			ANOVA							
11	6.9	2					5.398294	0.547164				df	SS	MS	F	Significance F		
12	7.0	0					5.390218	0.629977			Regression		2	387.8555	193.9278	202.0722	2.04E-32	
13	5.						5.384149	-1.35202			Residual		82	78.69502	0.959695			
14	6.4						5.380086	0.112449			Total		84	466.5505				
15	6.0						5.377981	-0.32888										
16	6.9	3					5.379938	0.610267				Coefficients	Standard Err.	t Stat	P-value	Lower 95%	Upper 95%	
17	6.1	2.5					5.383901	-0.24396			Intercept		6.589411803	0.31041	21.22811	4.59E-35	5.971908	7.206916
18	7.	2					5.389871	0.778618			x1		-2.914457293	1.484485	-1.96328	0.053004	-5.86757	0.038656
19	6.8	1.5					5.397848	0.448328			x2		10.03277678	1.467381	6.837202	1.32E-09	7.113689	12.95186
20	6.4	0.5					5.419821	-0.00722										
21	6.4	0					5.433818	-0.02859										
22	6.3	-0.5					5.449821	-0.10771										
23	5.7	-1					5.46783	-0.75517										
24	5.9	-1.5					5.487846	-0.49375										
25	6.2	-2					5.509869	-0.21582										
26	5.7	-2.5					5.533898	-0.81888										

- Plot of y (blue dots) and y_cap (red dots) seems to indicate that the model has captured the non-linear nature of y (how? why?)
- Plot of residuals also indicate no visible trend
- Model indicators like R2 and F-statistic also seem Ok
- Overall - the model seems to be good and acceptable based its performance on the train data
- We now need to check its performance on the test data

This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

The illustrated method of iteratively eliminating the **least important features** features, based on their p-values, is known as **BACKWARD FEATURES ELIMINATION**

Model performance on **Test Data**



- The y and ycap plots seem to indicate that the model, created using train data also performs reasonably well on the test data
- R2 value on test data seems Ok and close to the R2 value using train data
-

- The MSE values are differing, indicating some level of **overfitting** to the train data. However, the size of the test data is small, so the errors are possibly magnified.
- Usually, the technique of **cross-validation** is used - wherein multiple test data sets are used to evaluate the model. This results in an unbiased error estimate on test data.

OLS Regression Results

```
=====
Dep. Variable:                      y      R-squared:                 0.831
Model:                            OLS      Adj. R-squared:            0.827
Method:                           Least Squares      F-statistic:              202.1
Date:                Fri, 26 Jan 2024      Prob (F-statistic):        2.04e-32
Time:                     14:19:48      Log-Likelihood:            -117.33
No. Observations:                  85      AIC:                      240.7
Df Residuals:                      82      BIC:                      248.0
Df Model:                           2
Covariance Type:                nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	6.5894	0.310	21.228	0.000	5.972	7.207
x1	-2.9145	1.484	-1.963	0.053	-5.868	0.039
x2	10.0328	1.467	6.837	0.000	7.114	12.952

<=====

Omnibus:	1.380	Durbin-Watson:	2.318
Prob(Omnibus):	0.502	Jarque-Bera (JB):	1.164
Skew:	0.080	Prob(JB):	0.559
Kurtosis:	2.450	Cond. No.	23.2

<=====

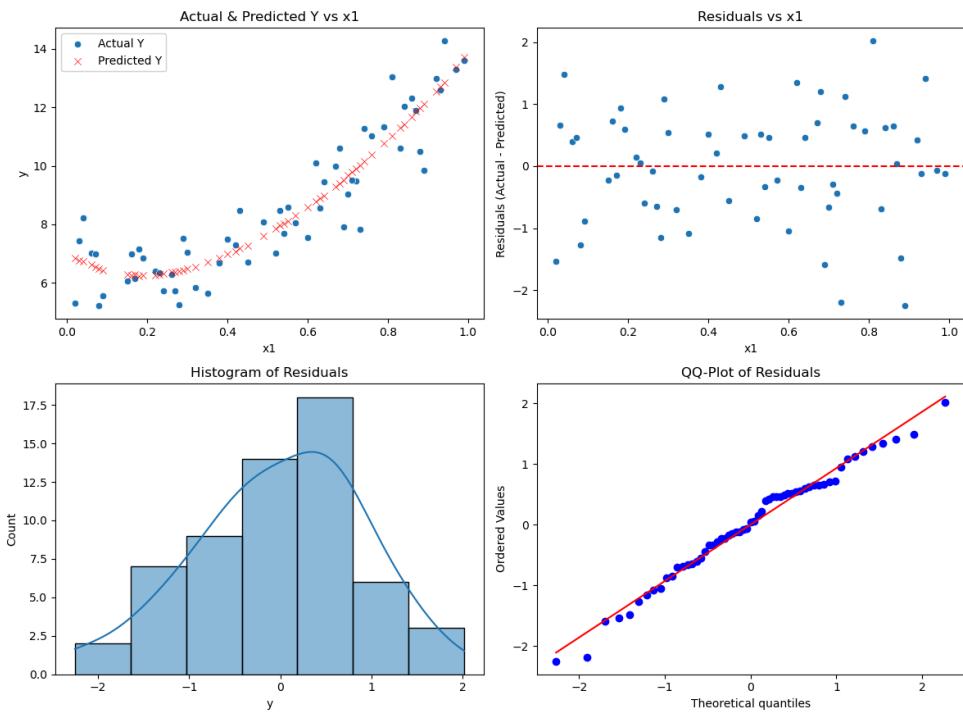
Exercise-2:

1. Re-create and validate the MLR model yourself, using the steps outlined in this document
2. Use the statsmodels based OLS function to repeat **all** these steps, and analyze the additional metrics created (Omnibus, Durbin-Watson, Jarque-Bera, AIC, BIC, Condition Number, etc.) at each stage

CREATE THE PLOTS & DATA SHOWN IN THE NEXT SLIDE

Note: In the data file '**data-set-for-MLR.xlsx**' the train and test data have been given on two different sheets

Sklearn - Train Analysis



DataFrame 'df' successfully loaded.

Target variable (y): y

Feature variables (x): ['x1', 'x2', 'x3', 'x4', 'x5']

First x variable for plotting: x1

Data Split: Train size = 59, Test size = 26

SCKIT-LEARN LINEAR REGRESSION

--- SKLEARN: Train Set Results ---

--- Sklearn Metrics ---

R2: 0.8602

MAE: 0.7375

RMSE: 0.9094

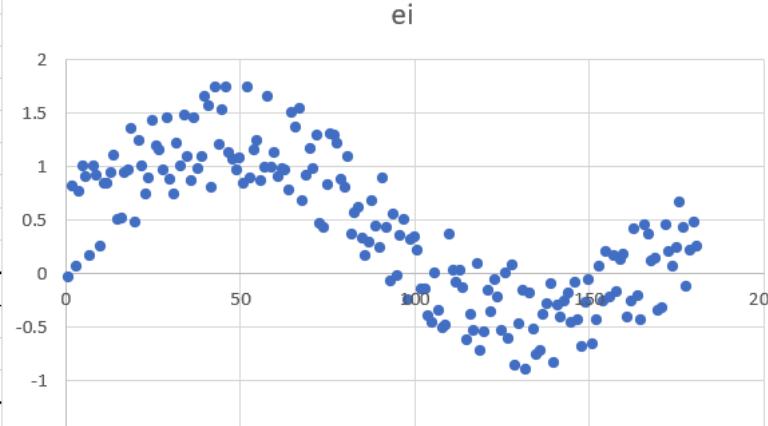
What if the relationship between the dependent variable (y) and the independent variables (X) is not linear as in the case below? We have seen that the regression errors are not random and the other important regression parameters like R² are also very low.

How do we remedy this situation?

	A	B	C	H	I	J	K	L	M	N	O	P	Q	R	S	T	U
1	y	x1		ycap	ei		SUMMARY OUTPUT										
2	0.038116834		0	0.078554	-0.04044												
3	0.896467788	0.005555556		0.083296	0.813172		Regression Statistics										
4	0.159545792	0.011111111		0.088032	0.071514		Multiple R	0.713818524									
5	0.863764416	0.016666667		0.092756	0.771008		R Square	0.509536885									
6	1.106349076	0.022222222		0.097463	1.008886		Adjusted R Square	0.506796867									
7	1.010169458	0.027777778		0.102147	0.908022		Standard Error	0.530607579									
8	0.278498289	0.033333333		0.106803	0.171695		Observations	181									
9	1.114230685	0.038888889		0.111424	1.002807												
10	1.029803908	0.044444444		0.116005	0.913799		ANOVA										
11	0.373869889		0.05	0.12054	0.25333			df	SS								
12	0.971634374	0.055555556		0.125024	0.84661		Regression	1	52.35633091								
13	0.975376766	0.061111111		0.129452	0.845925		Residual	179	50.3964481								
14	1.079774246	0.066666667		0.133817	0.945957		Total	180	102.752779								
15	1.242790434	0.072222222		0.138116	1.104675												
16	0.644698738	0.077777778		0.142341	0.502358		Coefficients	Standard Error									
17	0.656177067	0.083333333		0.146489	0.509688		Intercept	1.417122114	0.078553776	18.04015	3.95E-42	1.262112	1.572133	1.262112	1.572133	0.547807	
18	1.09549189	0.088888889		0.150554	0.944938		x1	-1.852833934	0.135870553	-13.6368	1.69E-29	-2.12095	-1.58472	-2.12095	-1.58472	12.74861	
19	1.115274736	0.094444444		0.154532	0.960743												
20	1.512547878		0.1	0.158416	1.354131												
21	0.639395626	0.105555556		0.162204	0.477192												

This file is meant for personal use by mepravintpatil@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.



We can see that the error plot is not random, and it follows a pattern. This indicates that forcing a **line** to model this data results in incorrect results. We need to **introduce non-linear independent variables** in the system so that the Multiple Linear Regression method can 'use' this non-linearity to produce the desired non-linear y_{cap} .

So, we introduce additional columns x_2, x_3, x_4 such that

$$x_2 = x_1 * x_1$$

$$x_3 = x_1 * x_1 * x_1$$

$$x_4 = x_1 * x_1 * x_1 * x_1$$

Note: The method is still Linear Regression. It is **Linear Regression of non-linear independent variables.**

The resulting regression method is known as **Polynomial Regression** - since polynomial terms are introduced as independent variable to handle non-linearity in y .

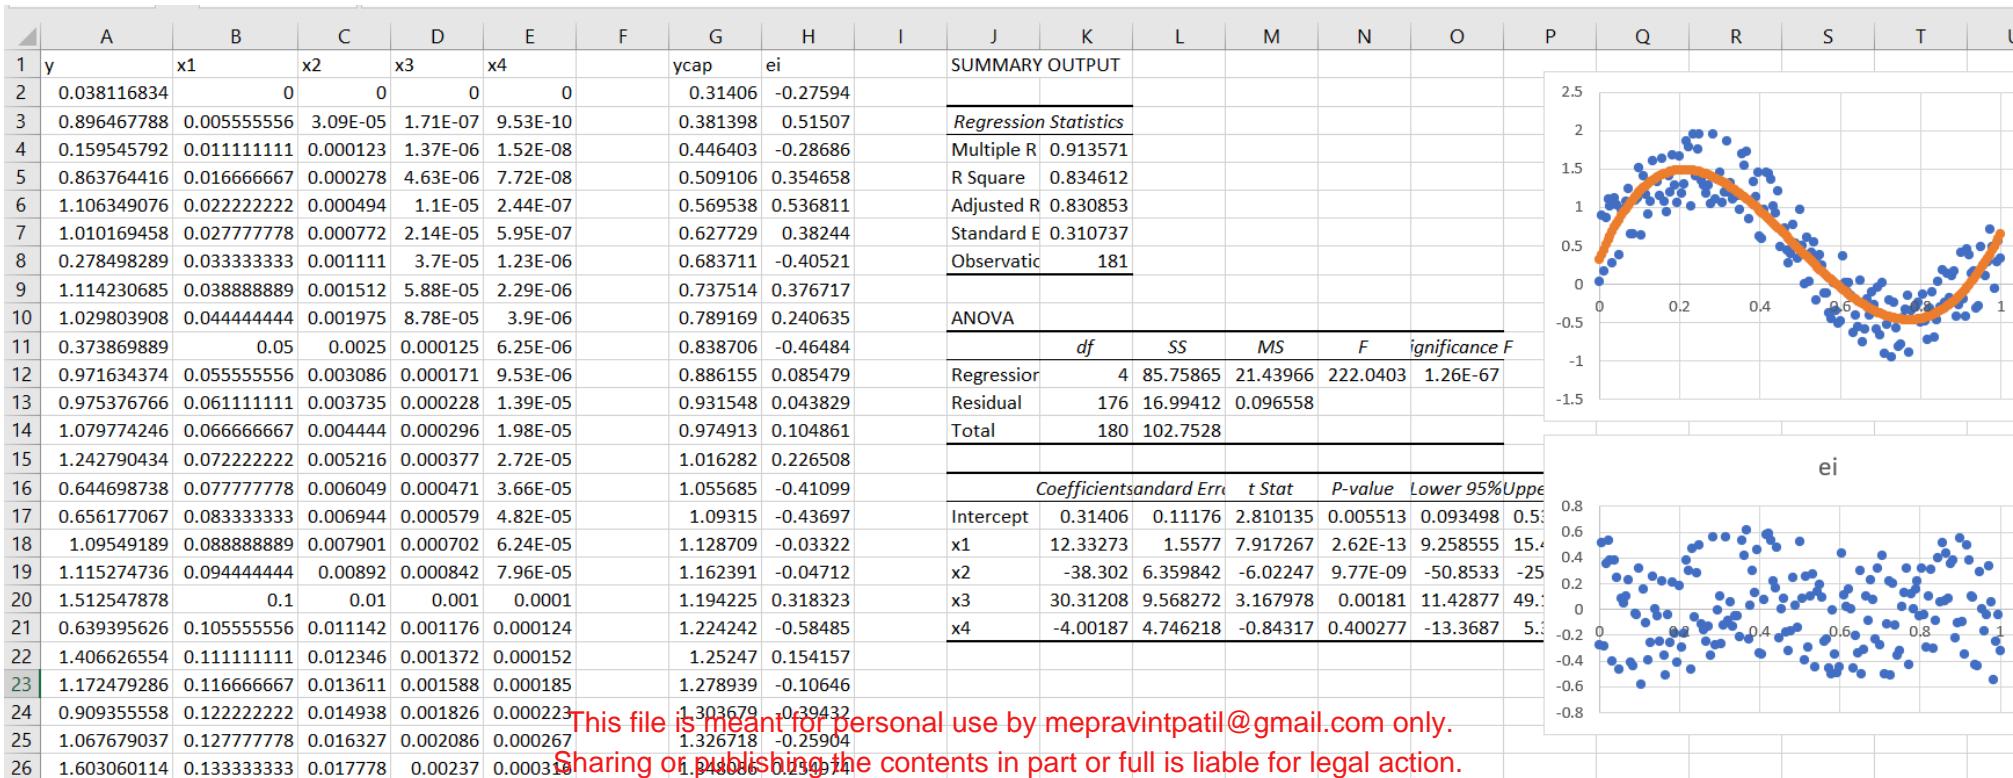
In general, introducing additional **x** variables to improve the performance of ML methods is known as **Feature Engineering**.

Hence, **Polynomial Regression** can be said to be an application of the **Feature Engineering** technique.

You are encouraged to create a dataset that is not good for being regressed by a line, but gets adequately represented by Polynomial Regression.

After introducing the polynomial terms and carrying out MLR, the results are as follows:

- The first chart shows y and y_cap (blue and orange)
- The second chart shows the error scatter plot
- We observe that the R² value is now quite good and all the p-values, except that for x₄, are much less than 0.05. This indicates that x₄ is not significant and needs to be dropped.



This file is meant for personal use by mepravintpatil@gmail.com only.
Sharing or publishing the contents in part or full is liable for legal action.

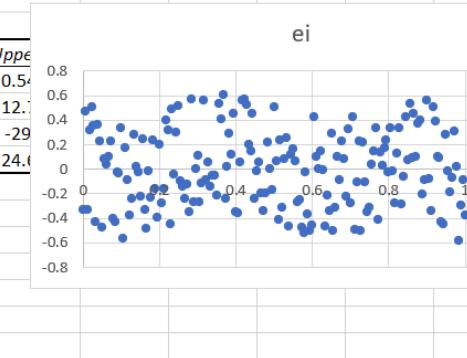
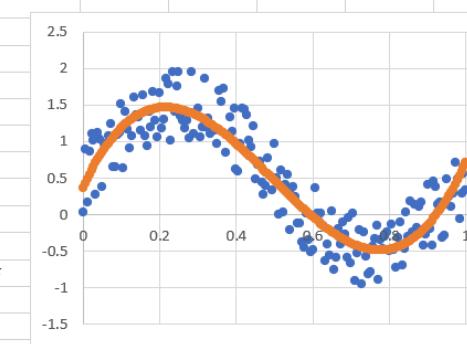
After dropping x4 from the model, the results are as follows:

- All p-values are now much lower than the threshold 0.05

This technique of starting off with all features and then **dropping** non-significant features one at a time is known as **Backward Feature Selection / Engineering.**

Backward feature engineering is a feature selection technique that removes features one by one until the model performance reaches a peak, and it is used to optimize the performance of the machine learning model by only including the most affecting feature and removing the least affecting feature.

y	x1	x2	x3	x4	ycap	ei	SUMMARY OUTPUT
0.038116834	0	0	0	0	0.369343	-0.33123	
0.896467788	0.005555556	3.09E-05	1.71E-07	9.53E-10	0.430539	0.465929	Regression Statistics
0.159545792	0.011111111	0.000123	1.37E-06	1.52E-08	0.48971	-0.33016	Multiple R 0.913205
0.863764416	0.016666667	0.000278	4.63E-06	7.72E-08	0.54688	0.316884	R Square 0.833943
1.106349076	0.022222222	0.000494	1.1E-05	2.44E-07	0.602072	0.504278	Adjusted R 0.831129
1.010169458	0.027777778	0.000772	2.14E-05	5.95E-07	0.655307	0.354862	Standard E 0.310483
0.278498289	0.033333333	0.001111	3.7E-05	1.23E-06	0.706611	-0.42811	Observatio 181
1.114230685	0.038888889	0.001512	5.88E-05	2.29E-06	0.756005	0.358226	
1.029803908	0.044444444	0.001975	8.78E-05	3.9E-06	0.803512	0.226292	
0.373869889	0.05	0.0025	0.000125	6.25E-06	0.849155	-0.47529	ANOVA
0.971634374	0.055555556	0.003086	0.000171	9.53E-06	0.892958	0.078676	df
0.975376766	0.061111111	0.003735	0.000228	1.39E-05	0.934943	0.040434	SS
1.079774246	0.066666667	0.004444	0.000296	1.98E-05	0.975133	0.104641	MS
1.242790434	0.072222222	0.005216	0.000377	2.72E-05	1.013552	0.229239	F
0.644698738	0.077777778	0.006049	0.000471	3.66E-05	1.050221	-0.40552	Significance F
0.656177067	0.083333333	0.006944	0.000579	4.82E-05	1.085164	-0.42899	Regressior
1.09549189	0.088888889	0.007901	0.000702	6.24E-05	1.118405	-0.02291	df 3
1.115274736	0.094444444	0.00892	0.000842	7.96E-05	1.149965	-0.03469	SS 85.69001
1.512547878	0.1	0.01	0.001	0.0001	1.179868	0.33268	MS 28.56334
0.639395626	0.105555556	0.011142	0.001176	0.000124	1.208137	-0.56874	F 296.3007
1.406626554	0.111111111	0.012346	0.001372	0.000152	1.234795	0.171832	Significance F 9.58E-69
1.172479286	0.116666667	0.013611	0.001588	0.000185	1.259864	-0.08738	Total 180
0.909355558	0.122222222	0.014938	0.001826	0.000223	1.283368	-0.37401	
1.067679037	0.127777778	0.016327	0.002086	0.000267	1.30533	-0.23765	
1.603060114	0.133333333	0.017778	0.00237	0.000316	1.344718	0.023186	
1.367903685	0.138888889	0.01929	0.002679	0.000372			



Backward v/s Forward Feature Engineering

Forward Feature Engineering	Backward Feature Engineering
Starts with an empty feature set and iteratively adds one feature at a time based on their performance	Starts with a complete set of features and removes features one by one until the model performance reaches a peak
Goal is to identify the most accurate and informative features that contribute to the predictive power of the model	Goal is to identify the most accurate and relevant features that can be used in a model
Iteratively adds features to the model	Iteratively removes features from the model
Can be a more time-consuming process than backward feature engineering	Can be a more systematic approach than forward feature engineering
Can be useful when the number of features is relatively small	Can be useful when the number of features is relatively large
Can be prone to overfitting if too many features are added to the model	Can be prone to underfitting if too many features are removed from the model
Can be used in combination with backward feature engineering to optimize the feature selection process	Can be used in combination with forward feature engineering to optimize the feature selection process

In summary, forward feature engineering and backward feature engineering are two techniques used in machine learning for selecting relevant features to include in a model. Forward feature engineering starts with an empty feature set and iteratively adds one feature at a time based on their performance, while backward feature engineering starts with a complete set of features and removes features one by one until the model performance reaches a peak. Both techniques have their advantages and disadvantages and can be used in combination to optimize the feature selection process.

Exercise-3

- Try Backward Feature Elimination by adding polynomial and other relevant functions as base features to the data set in **non-linear-data-set-for-regression.csv**
- Try the Forward Feature Selection method for the same dataset
- Try the mixed approach (forward + backward) feature selection on the dataset.

Exercise-4

- Perform Linear Regression by adding appropriate features (polynomial / others) to the uploaded dataset **sine-segment-perturbed.csv**. What conclusions can you make?